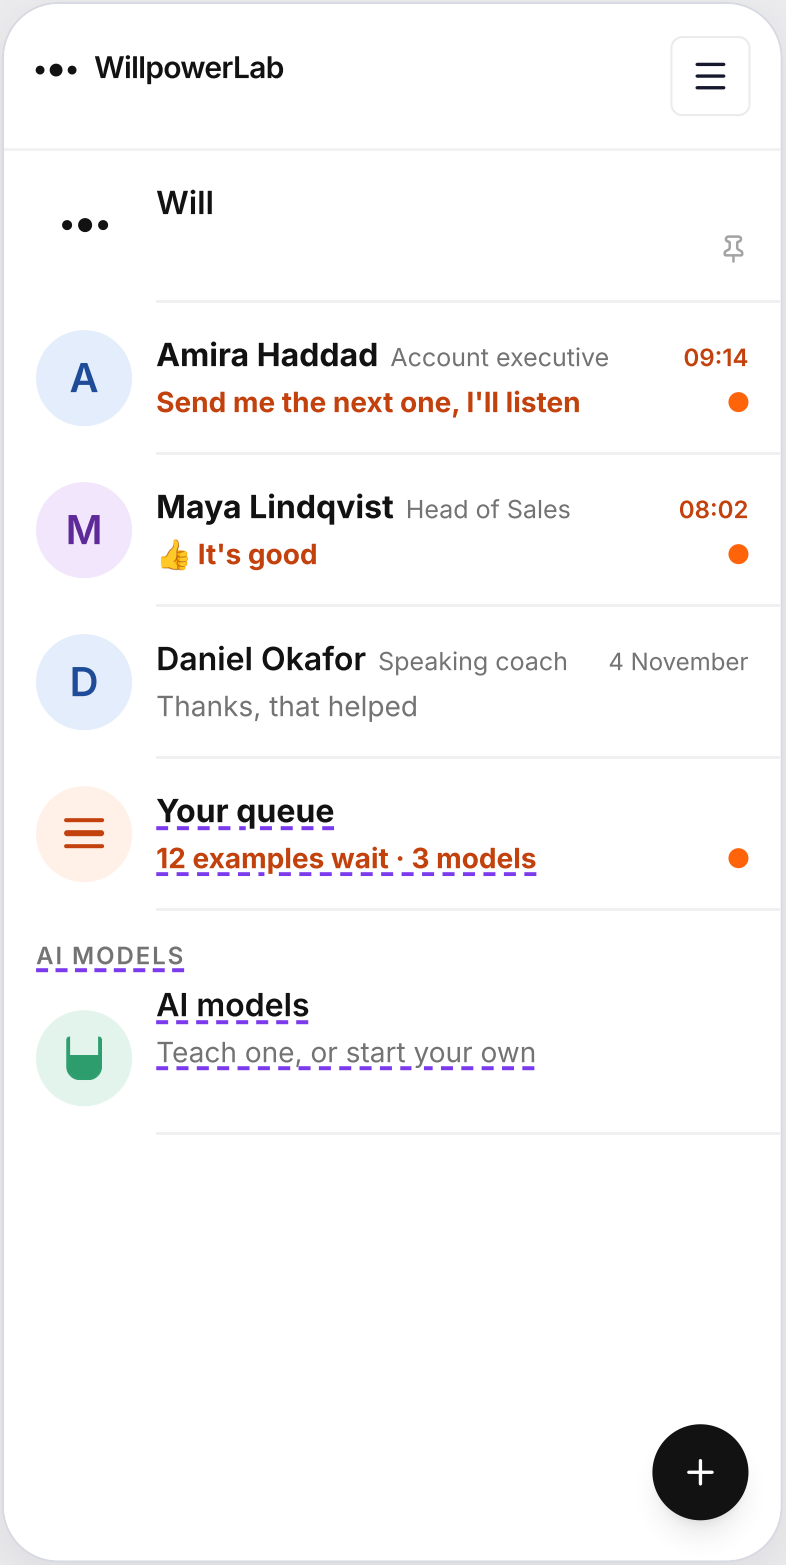
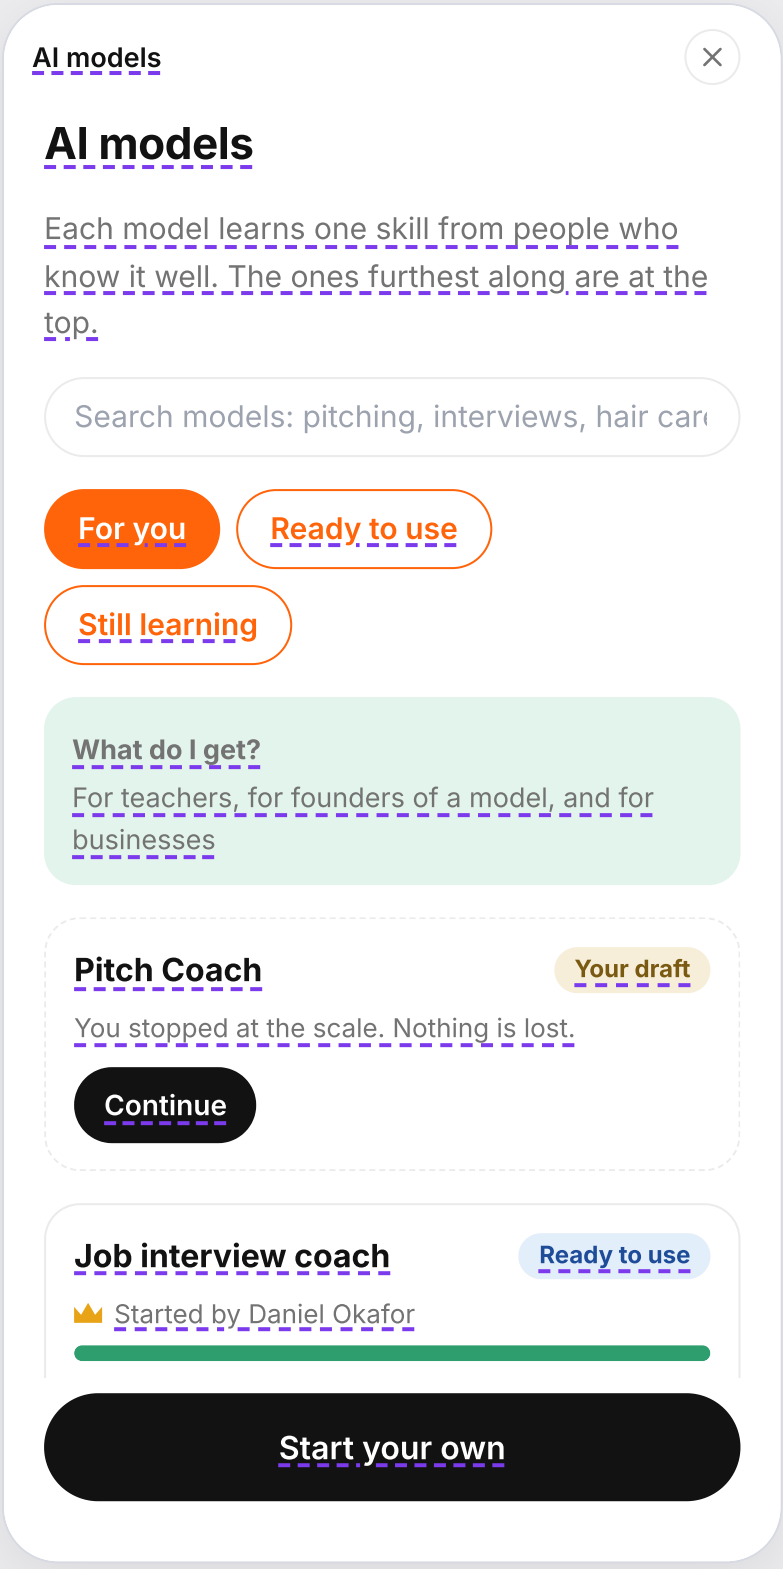
The model launchpad

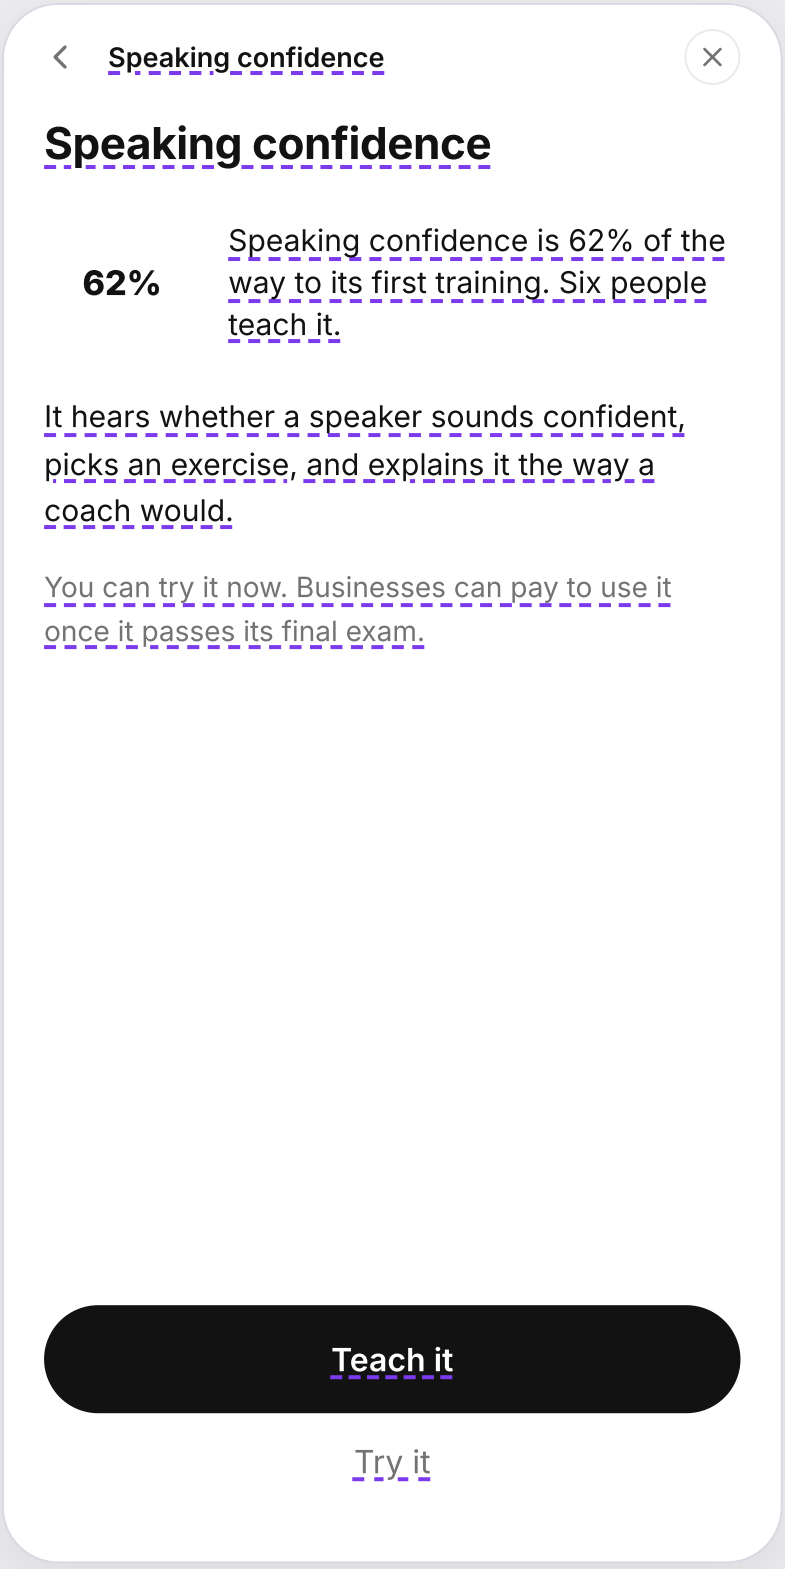
A clickable prototype of N528 on the app's real screens: finding models under your chats, teaching one, starting your own in single steps, starting your own (with your existing data, if you have some), using a ready one in your own business through its API, and downloading its data. Tap inside the phone, or use Back and Next. Sample names and numbers. Nothing here is built or locked.



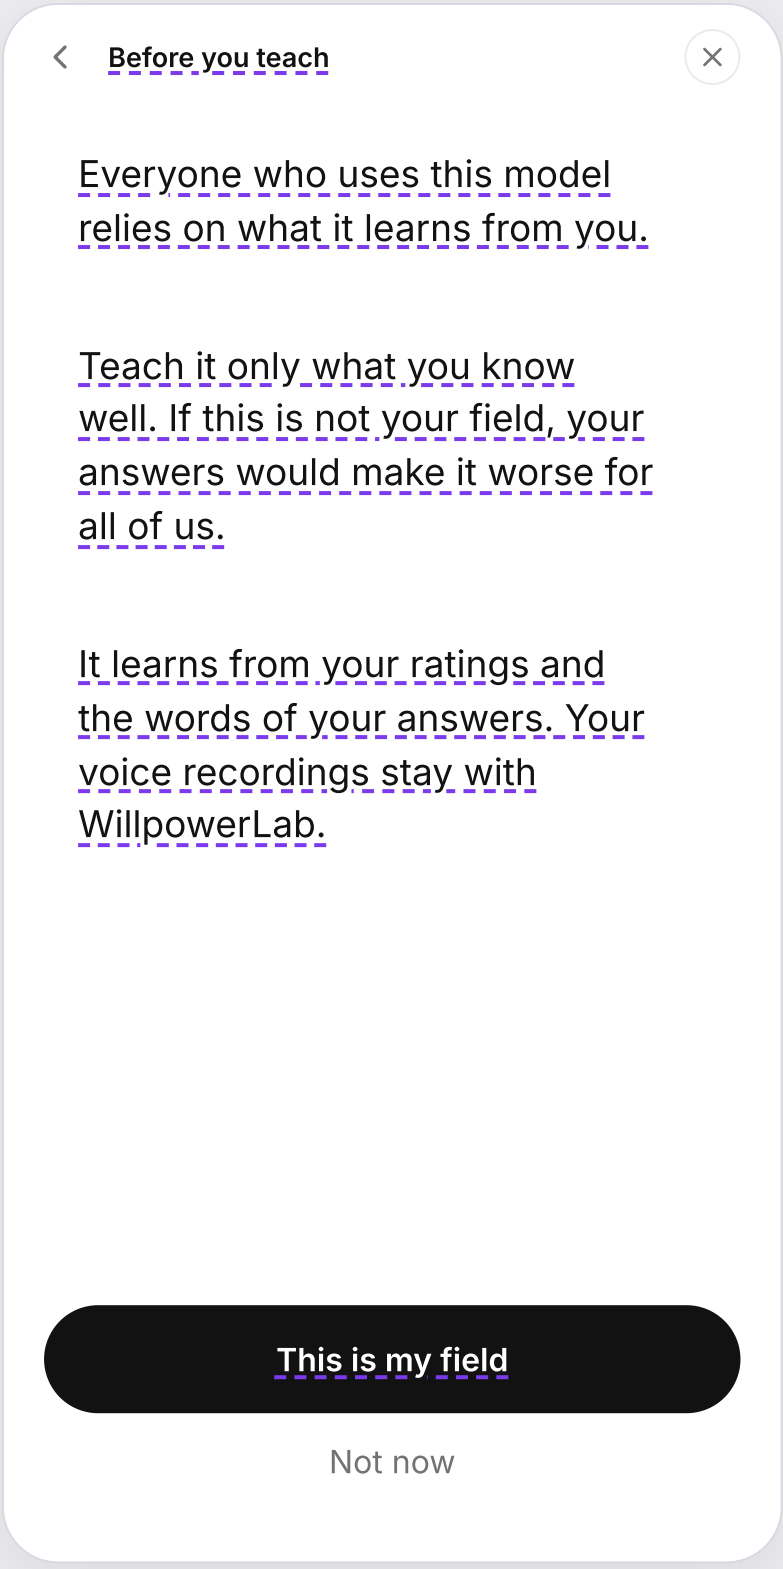
1/8 Your list of chats. Under the people, one new row: **AI models**. Tap it.



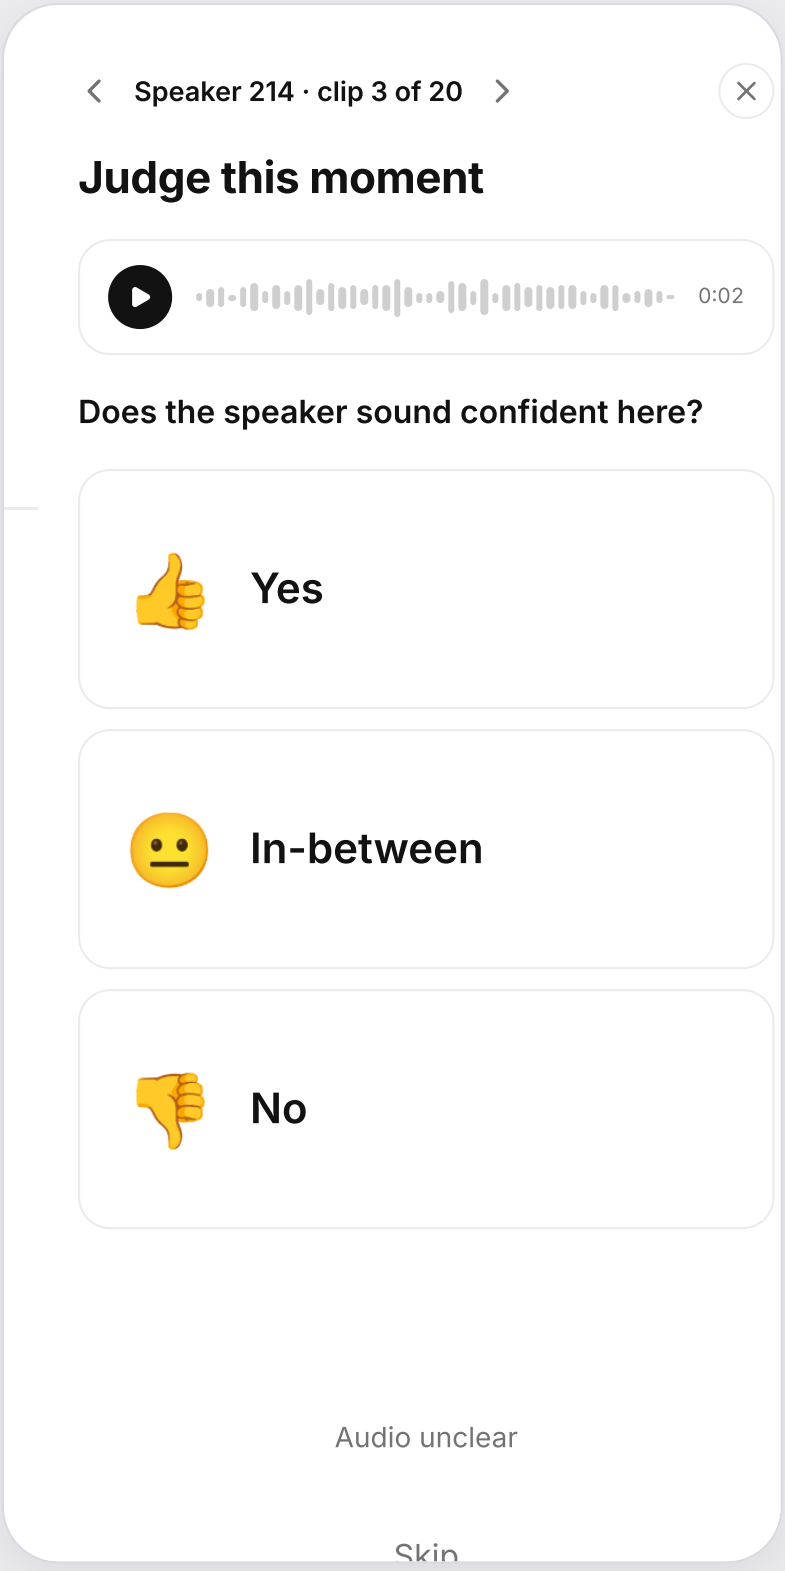
2/8 The list. The most mature first: one is ready to use, the others show how far they are from their first training, and who started each (the small crown). Tap **What do I get?**, **Use it**, **Teach it**, or **Start your own**.



3/8 A public model's page, as anyone sees it.



4/8 The disclaimer at the door. Anyone may come in; the next screens check quietly.



5/8 Teaching looks like the coach panel's Judge. A few of the first 20 clips already have a known answer; that is the quiet trust check. Tap an answer.

Speaking confidence

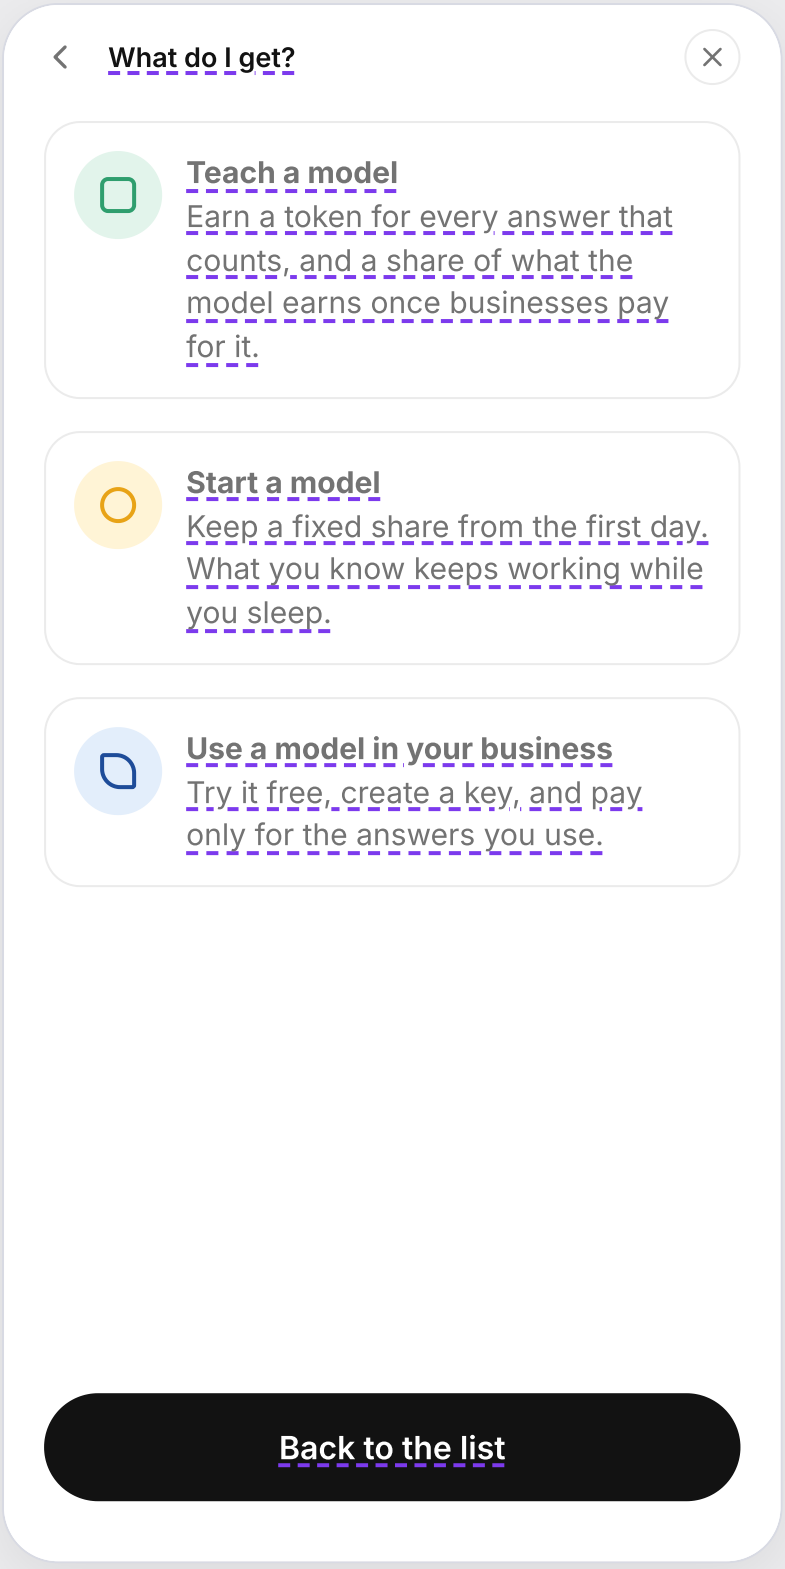
100% Your answers now count.

You earn a token for each answer that counts. When businesses start paying to use this model, you also get a share of what it earns.

Keep teaching

Back to the list

6/8 After the first 20 answers, if they hear what the experts hear.



7/8 What each person gets, in plain words: teachers, founders of a model, and businesses.

AI models

AI models

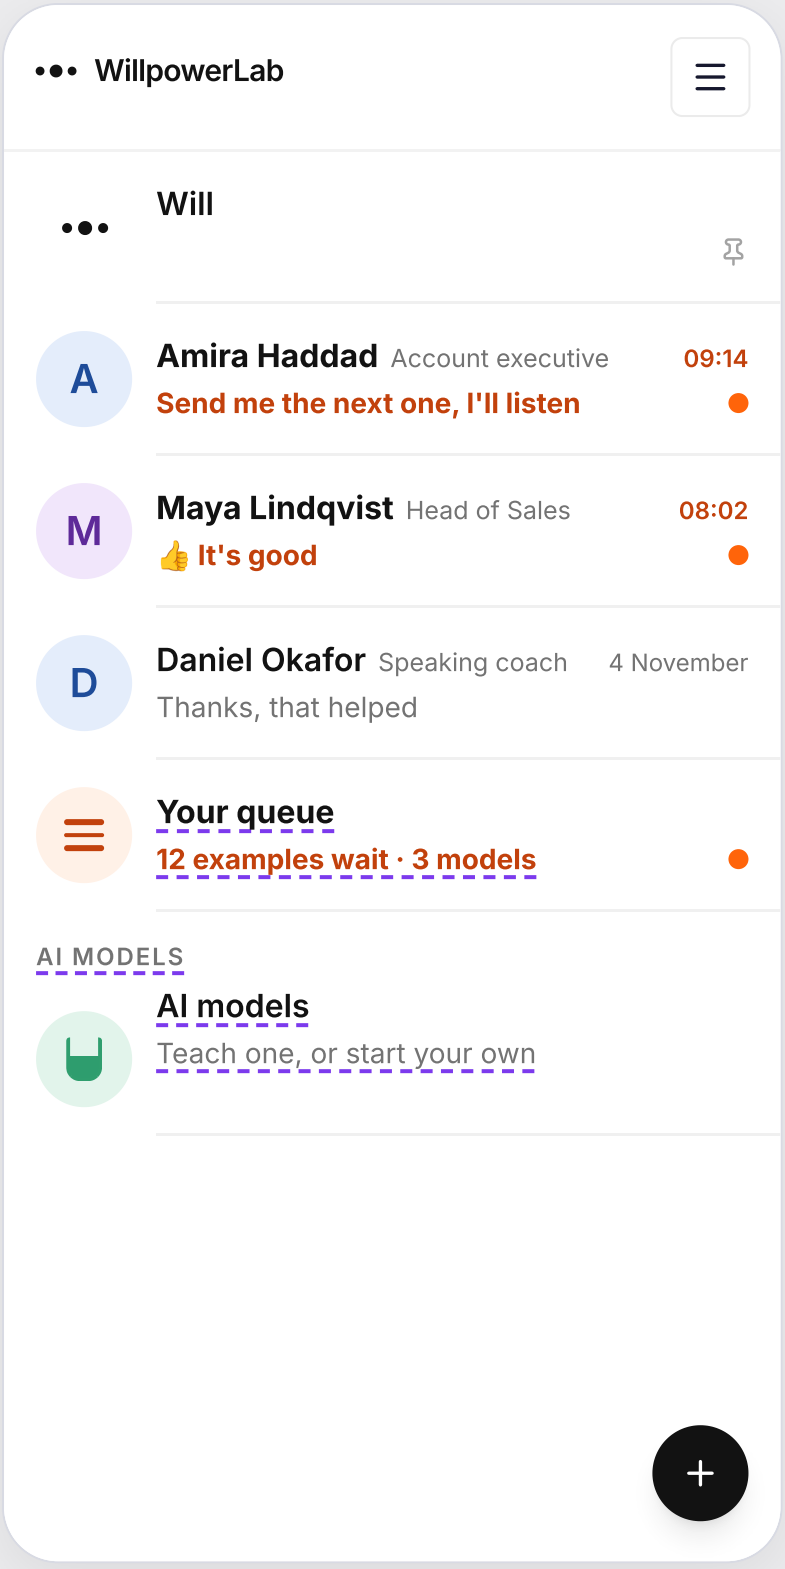
No models here yet.

Nobody has started a model about this.
You could be the first.

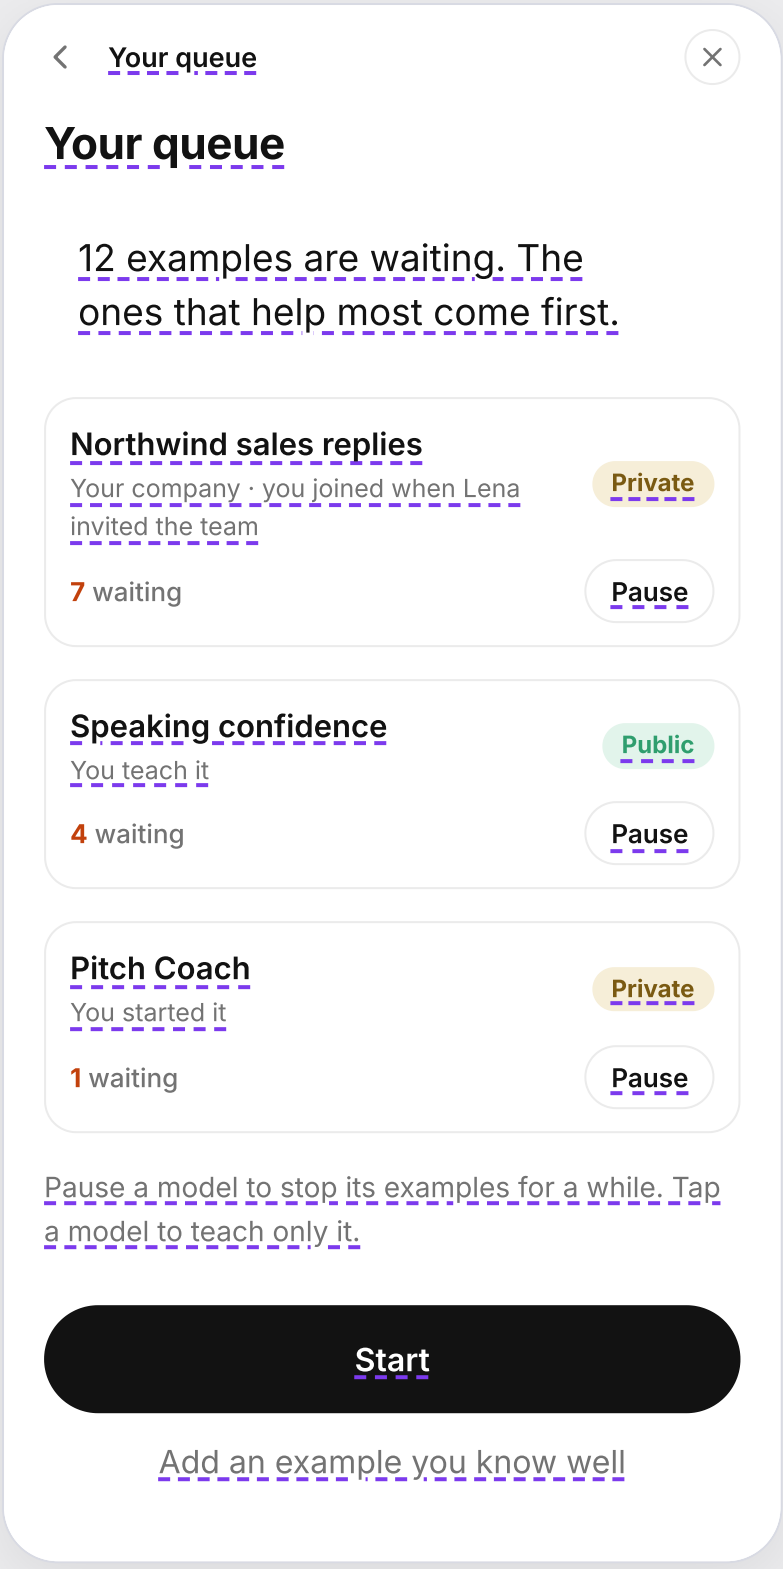
Start the first one

Clear the search

8/8 New, from the UX review: the list when a search or a filter finds nothing.



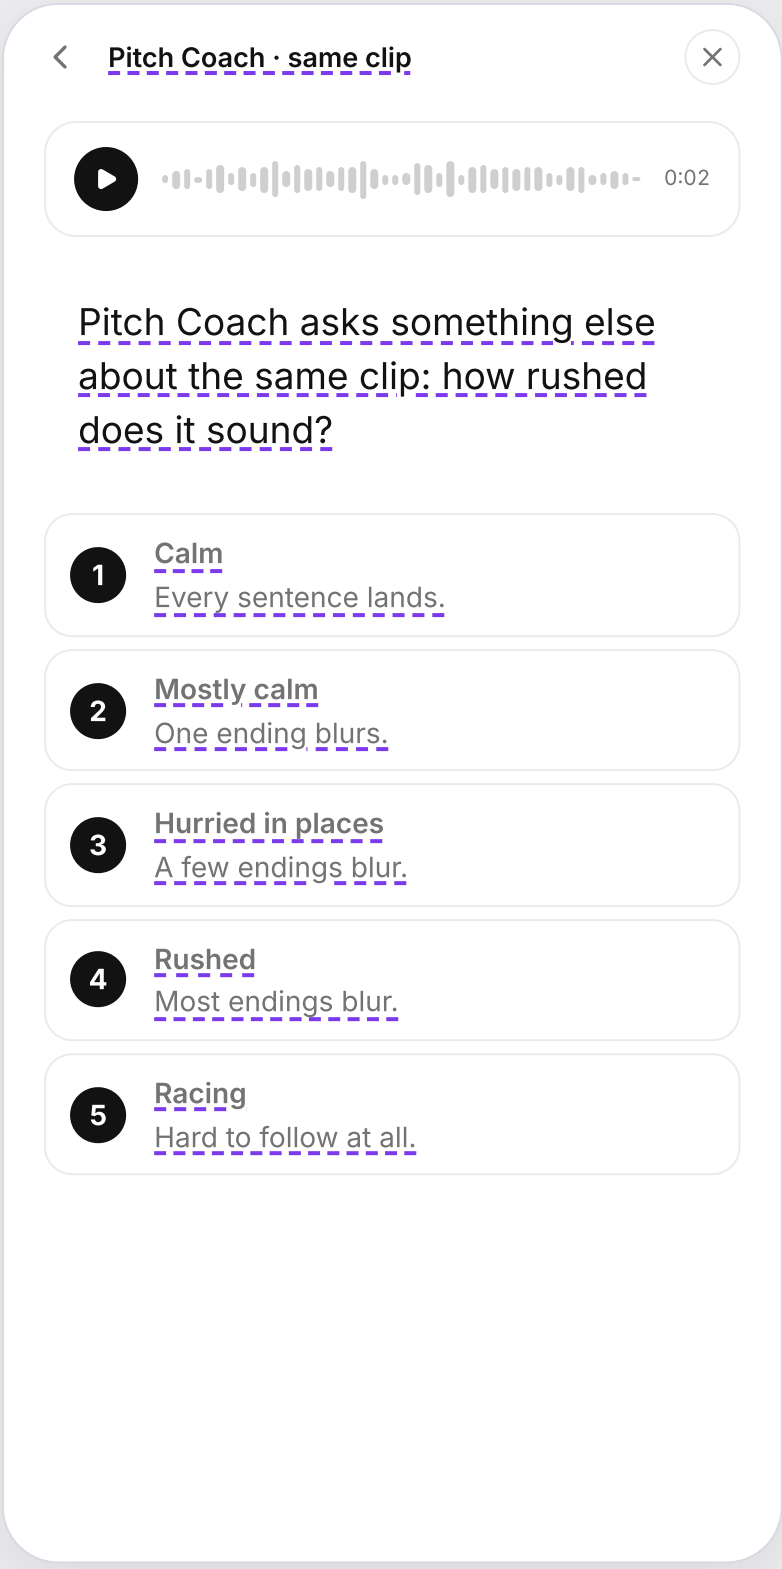
1/9 Your list of chats. The coach's row from N521 is now **Your queue**: everything you teach, in one place. Tap it.



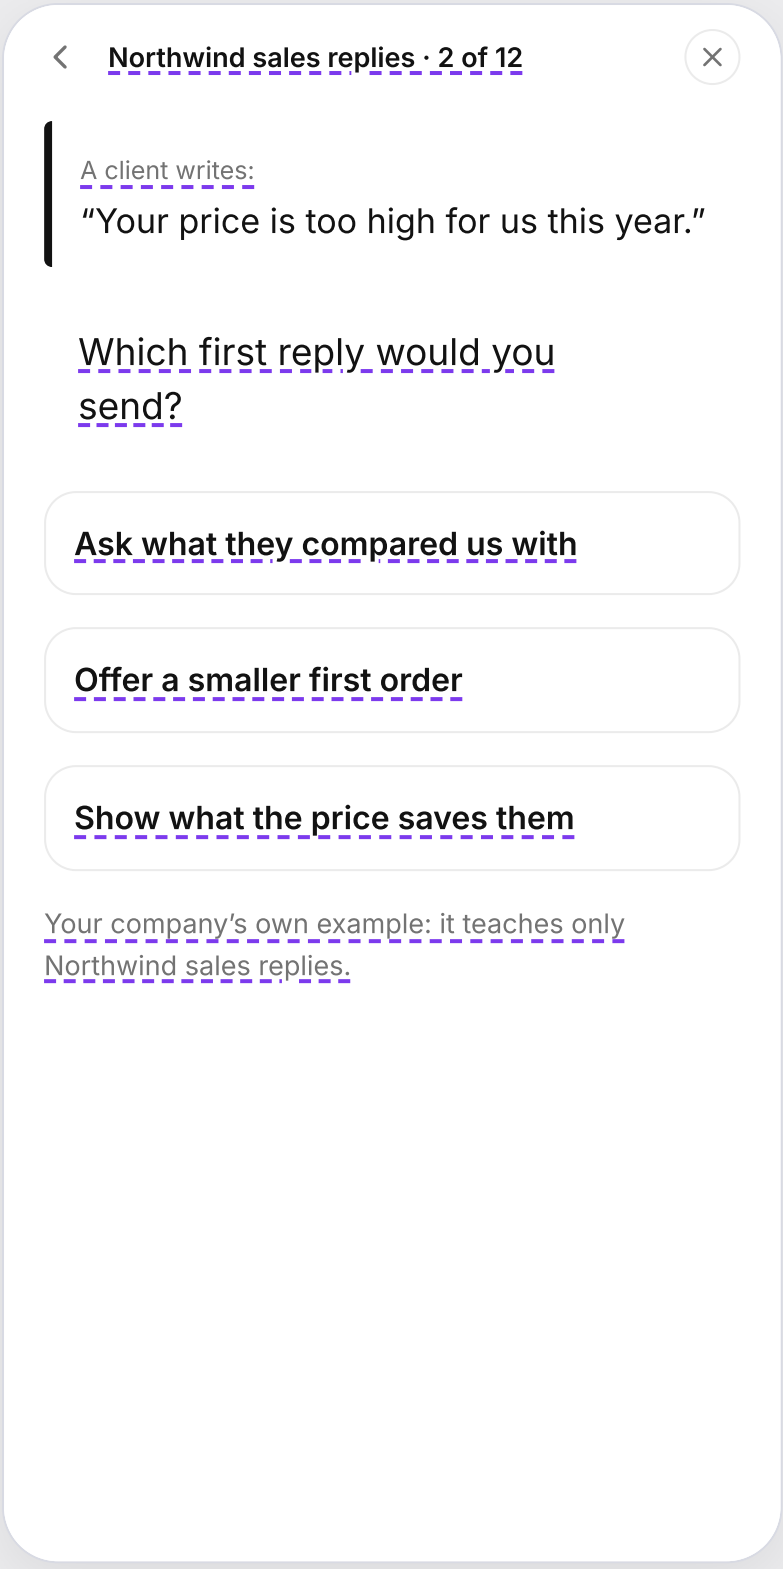
2/9 One queue across every model you teach: your company's private one and the public ones. The system orders it so each model gets what helps it most; you can pause any of them.

Speaking confidence · 1 of 12

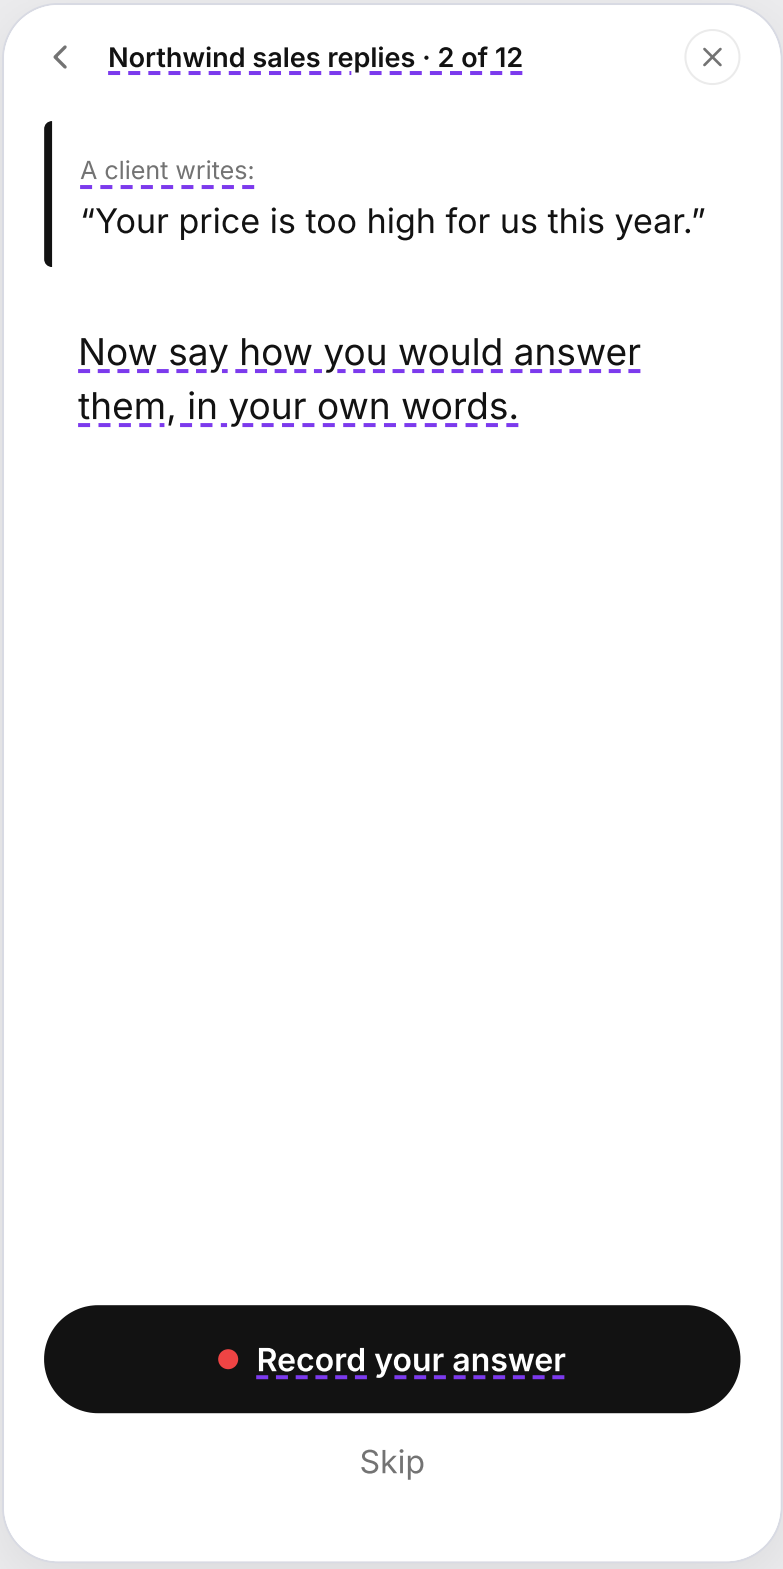
3/9 Fixed, from the data-science review: one answer counts for two models only when both ask the same question on the same scale, as Speaking confidence and Interview voice do here. The clip comes from a public source, WillpowerLab speakers who said yes. Tap an answer.



4/9 Pitch Coach asks a different question about the same clip, on its own 1 to 5 scale, so it comes right after, without loading the clip again. Tap a point.



5/9 A text example for your company’s private model. Its recipe card asks for a pick from a list, then your own answer, so the screen follows the card. Private data never feeds a public model. Tap a reply.



6/9 The answer is spoken, as every answer in WillpowerLab is: the app’s black button with the red dot.

Add an example

Do you know a case this model should learn from? Add it, then answer it yourself.

Record one
Say the case in your own voice

Upload a file
A recording, a photo or a written case

It counts once another teacher agrees with your answer.

7/9 Teachers can add examples they know well. Another teacher has to agree before it counts, so nobody can teach a model alone.

Pitch Coach

Version 2 is ready

100% Version 2 of Pitch Coach passed its final exam.

On the 100 exam cases:

14	3	83
version 2 better	version 1 better	the same

Version 2 is clearly better: a gap of 14 against 3 is very unlikely to be luck.

Version 2 agrees with the experts	**82 times in 100**
Two experts agree with each other	**84 times in 100**

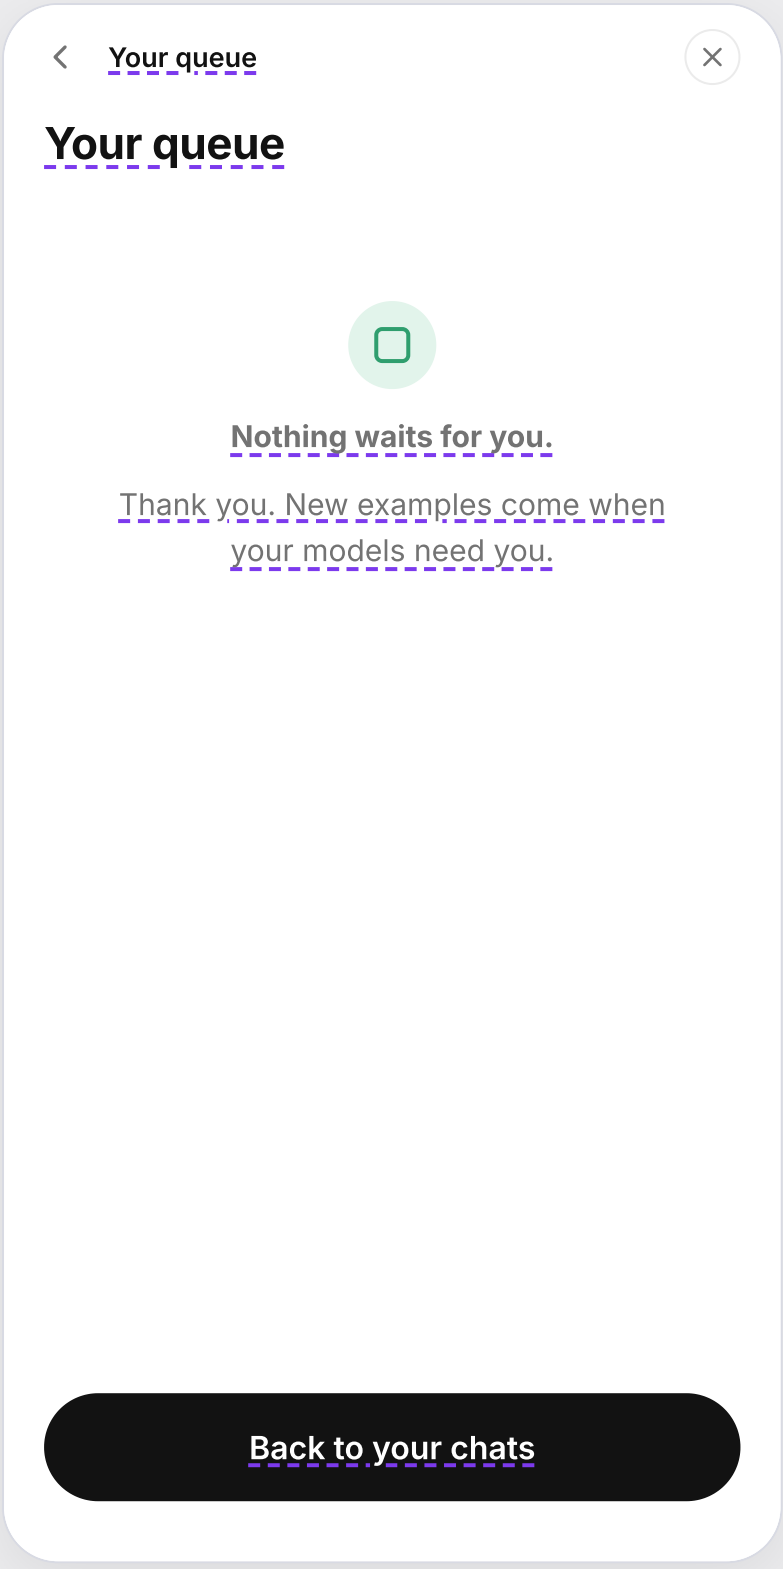
Had it been closer, say 9 against 6, you would see “Not clearly better yet”, and no Release button.

Releasing it is your choice. Businesses on version 1 keep it until they decide to switch.

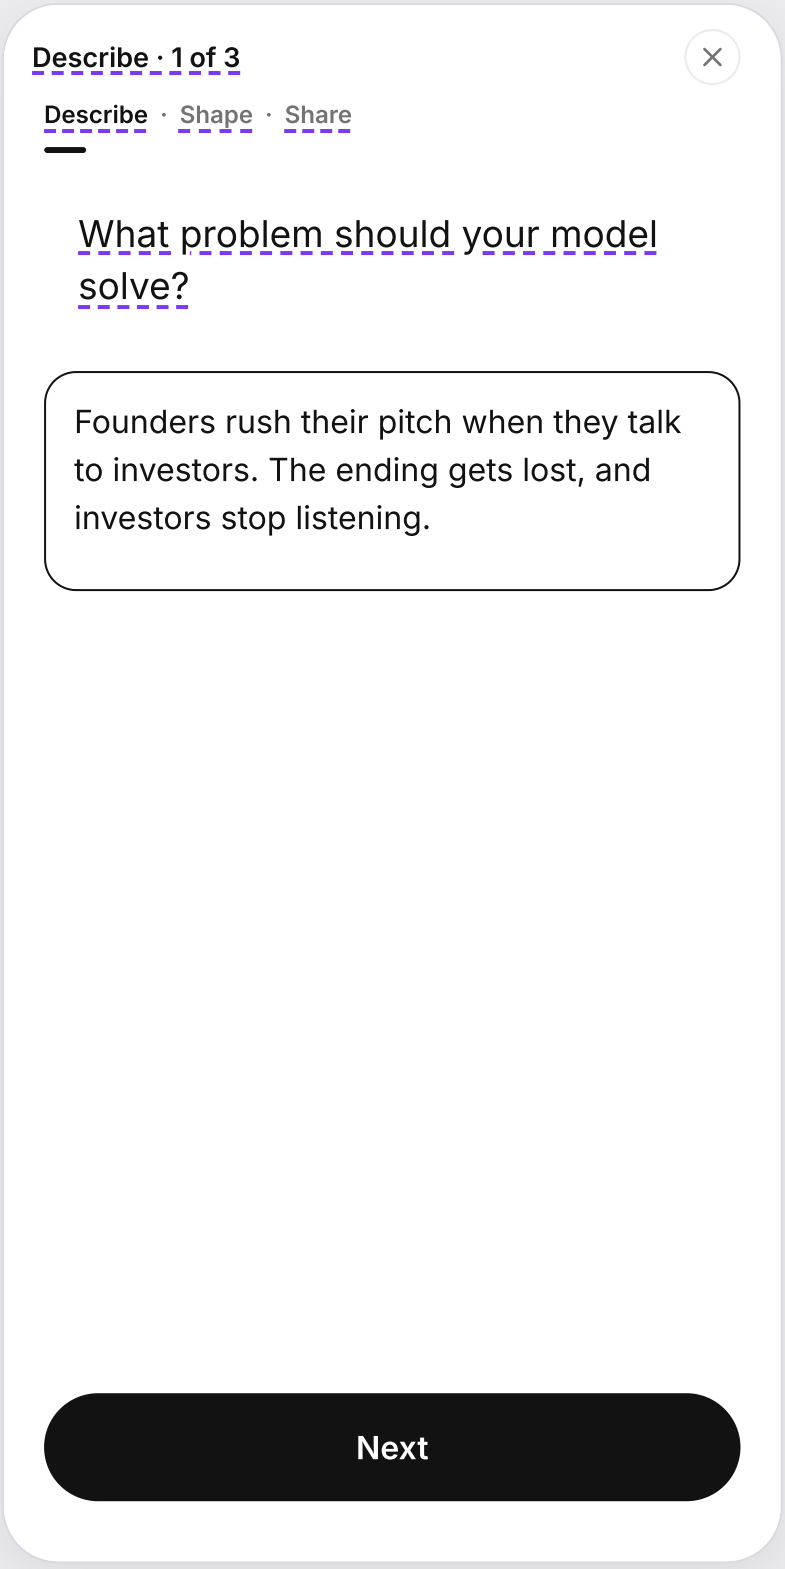
Release version 2

Not yet

8/9 Fixed, from the data-science review: versions are compared case by case on the same exam, and Release appears only when the difference is real, not luck. The founder makes one of their four decisions. (Sample numbers.)



9/9 New, from the UX review: the queue when it is empty.



1/18 Starting a model. **Changed, from the UX review:** the setup has three chapters, Describe · Shape · Share, with a thin bar that fills as you go; optional steps say so. × keeps your draft; the list shows it with Continue. Three short answers first, one per screen.

Describe · 2 of 3

Describe · Shape · Share

How do you solve it today?

I listen to their pitch, tell them where they sped up, and give them one exercise for the ending.

Next

2/18 The second answer.

Describe · 3 of 3

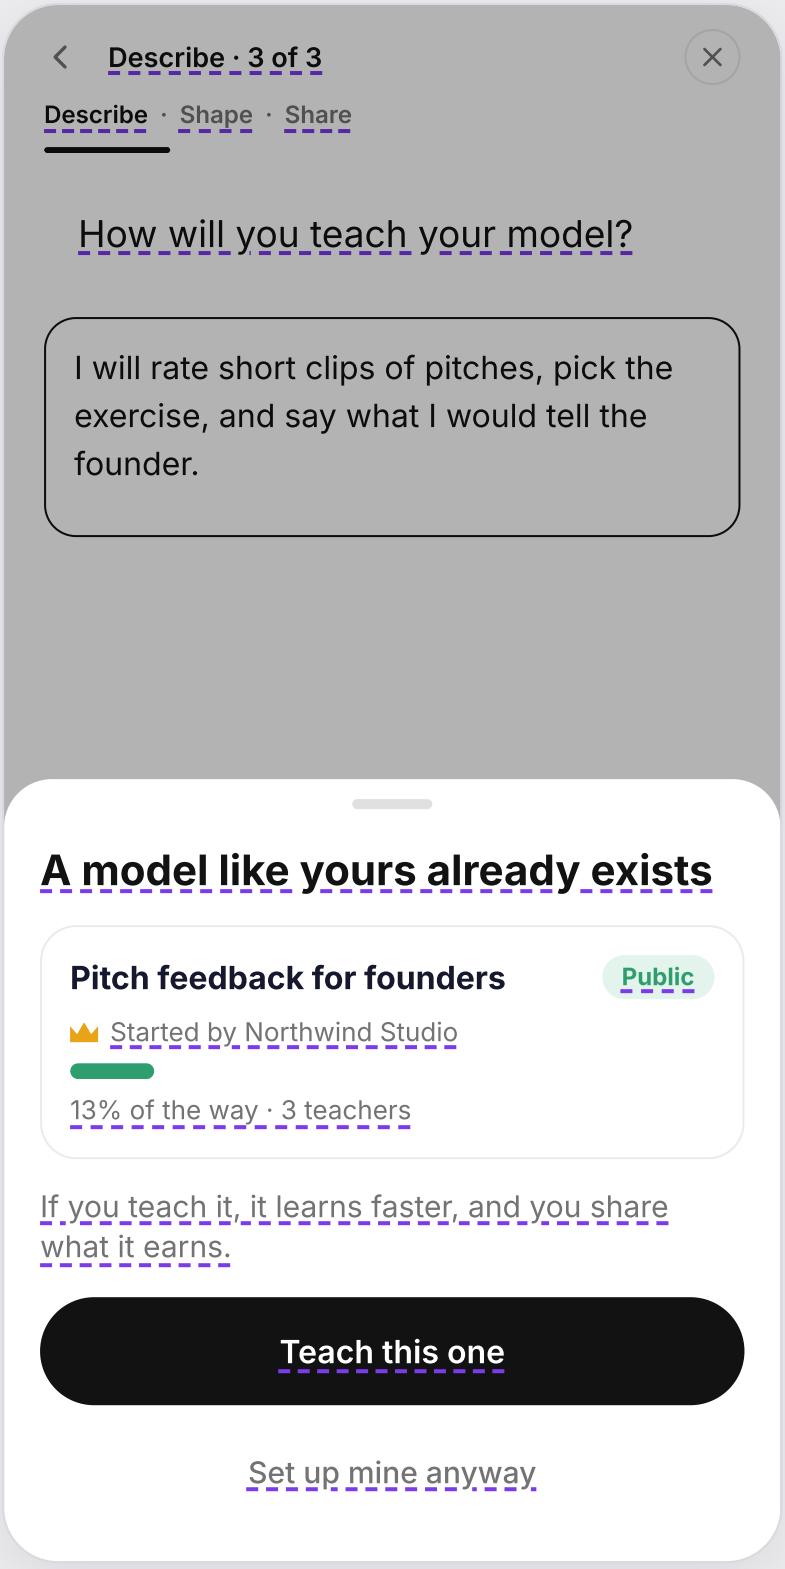
Describe · Shape · Share

How will you teach your model?

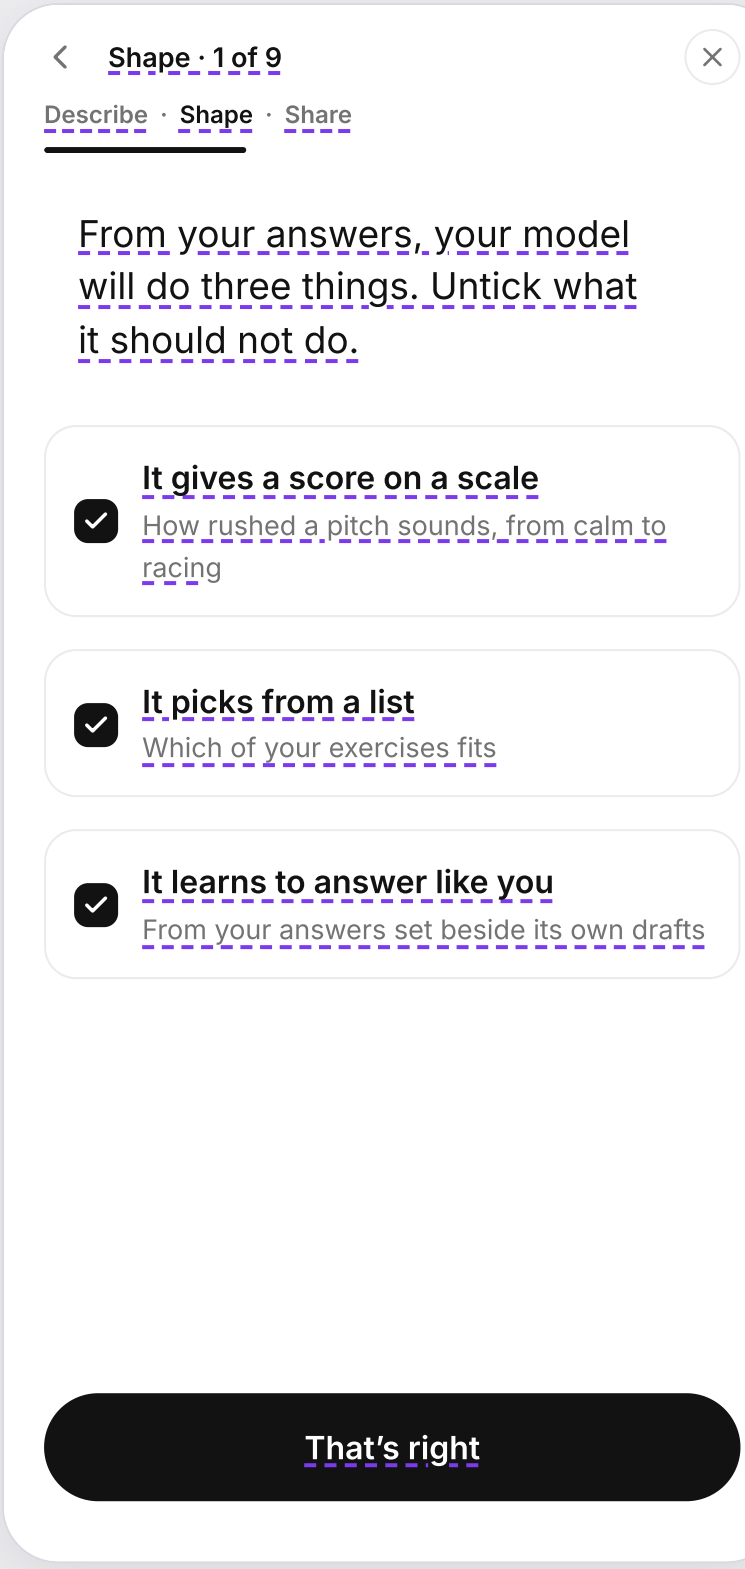
I will rate short clips of pitches, pick the exercise, and say what I would tell the founder.

Next

3/18 The third answer.



4/18 Changed, from the UX review: "A model like yours already exists" slides up over the third answer as a sheet, in the app's own sheet style, instead of being a numbered step that only sometimes appears.



5/18 The app proposes the kind of model: a prediction on a scale, a choice from a list, or learning to answer like you. One tap to accept.

Shape · 2 of 9

Describe · Shape · Share

One example is a pitch clip of about 30 seconds. Your teachers will rate each clip on this scale.

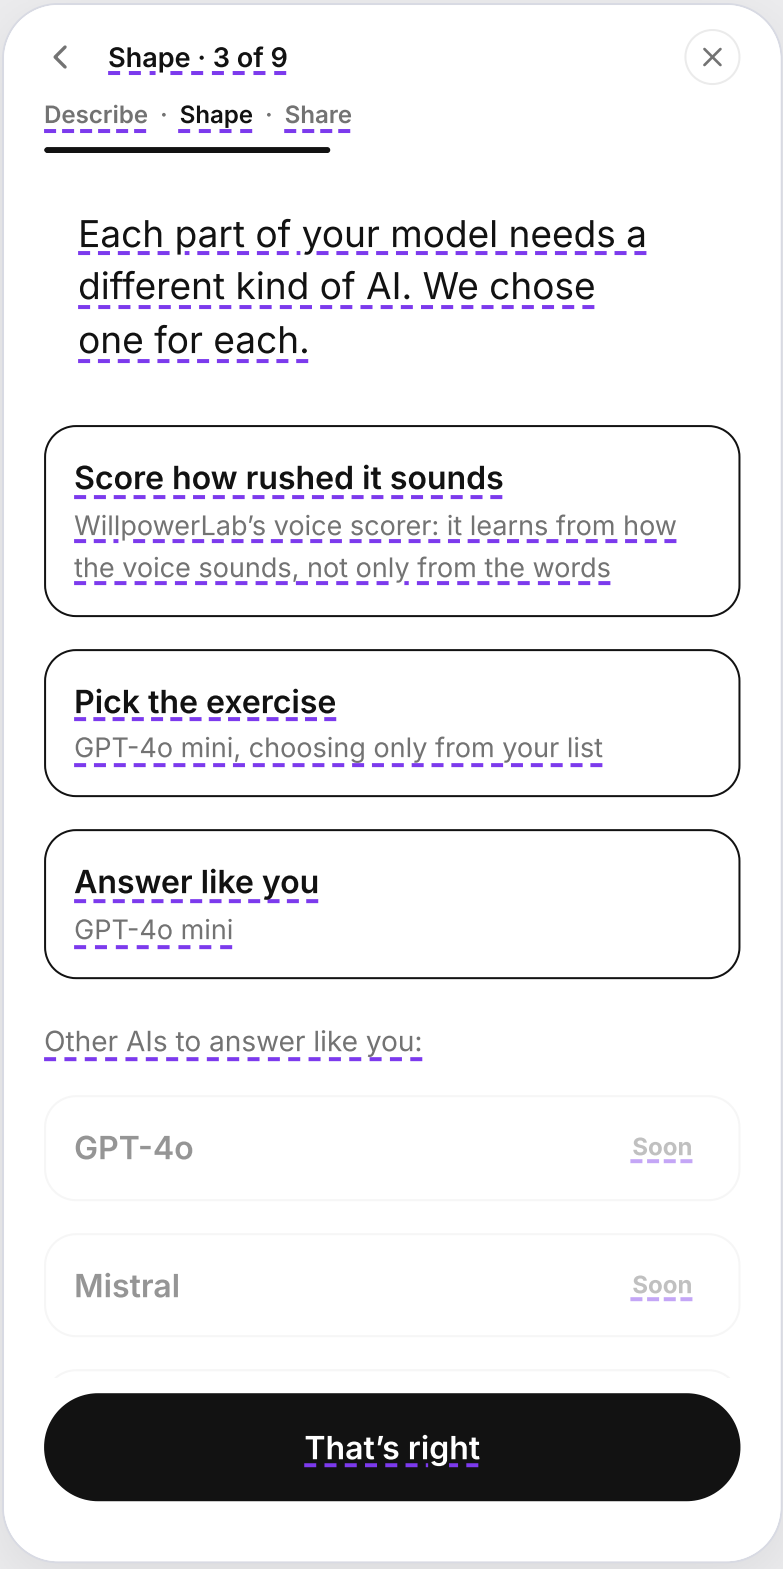
1. **Calm**
   Every sentence lands.
2. **Mostly calm**
   One ending blurs.
3. **Hurried in places**
   A few endings blur.
4. **Rushed**
   Most endings blur.
5. **Racing**
   Hard to follow at all.

Clear words for each point help your teachers agree. When they agree, the model learns faster.

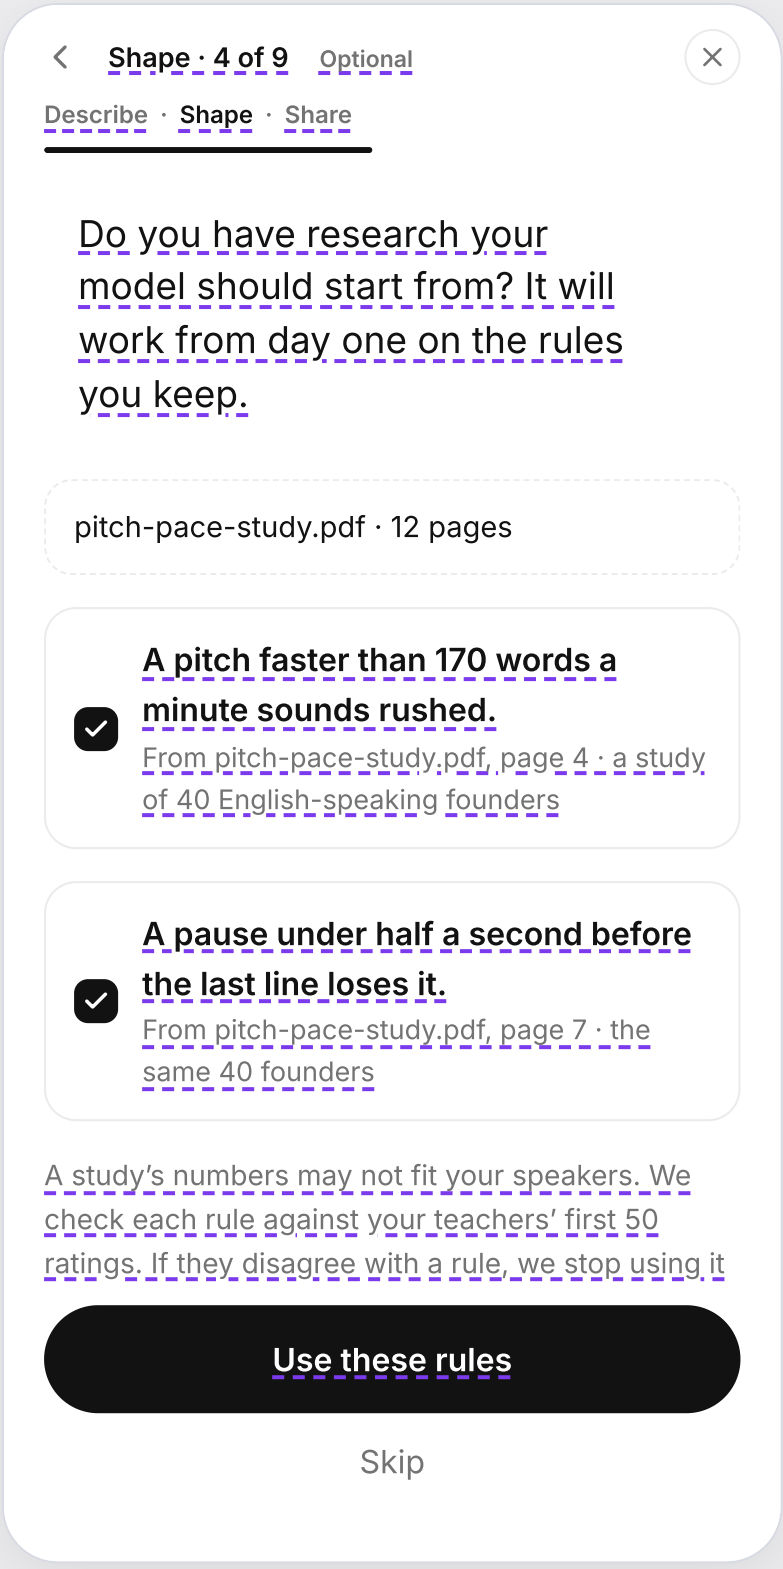
Use this scale

Start with this, finish the rest later

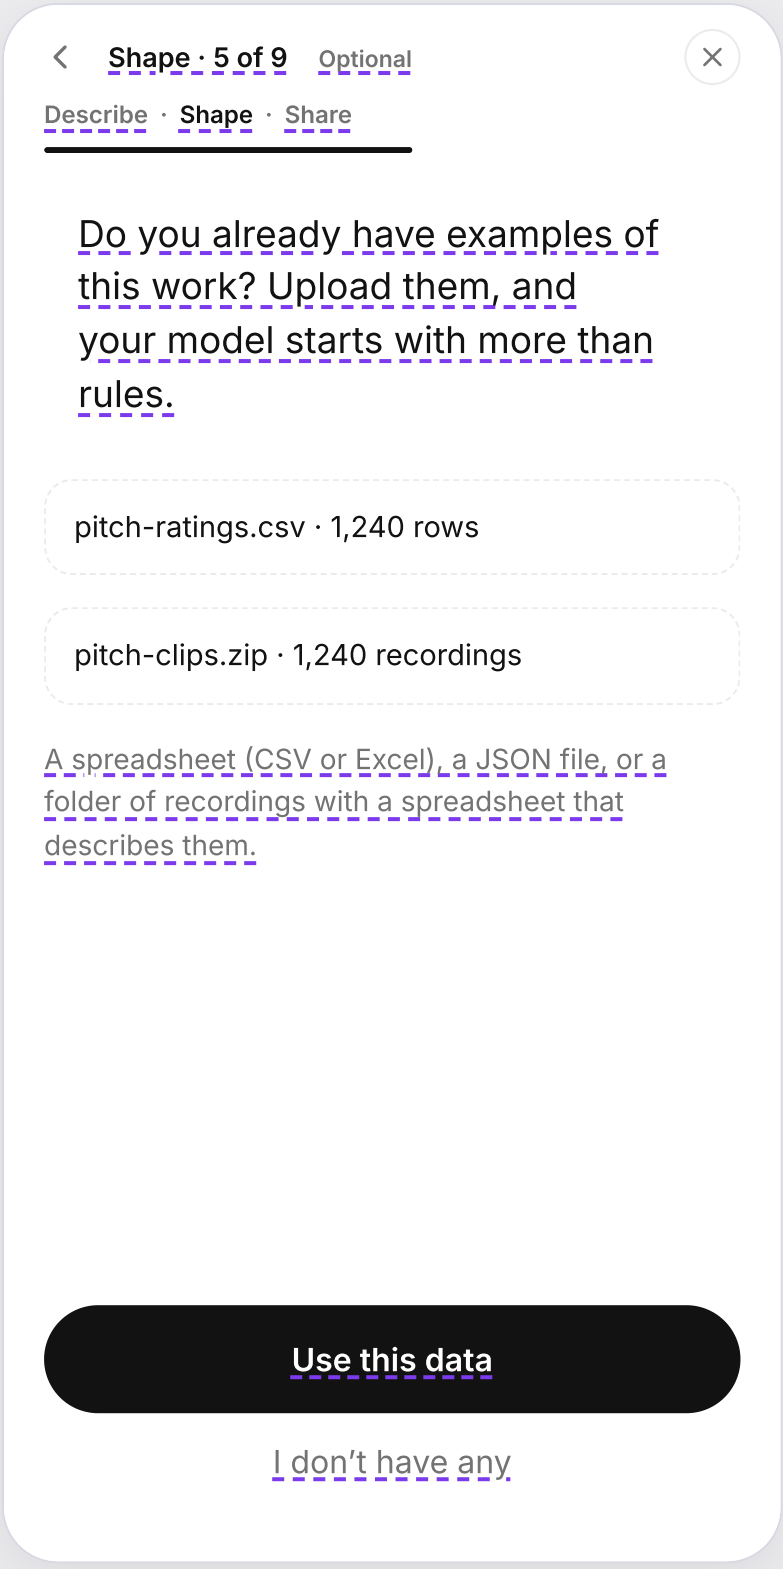
6/18 New, from the data-science review. What one example is, and a scale with words for every point. Teachers who share clear anchors agree far more often, and a model can only be as good as its teachers' agreement.



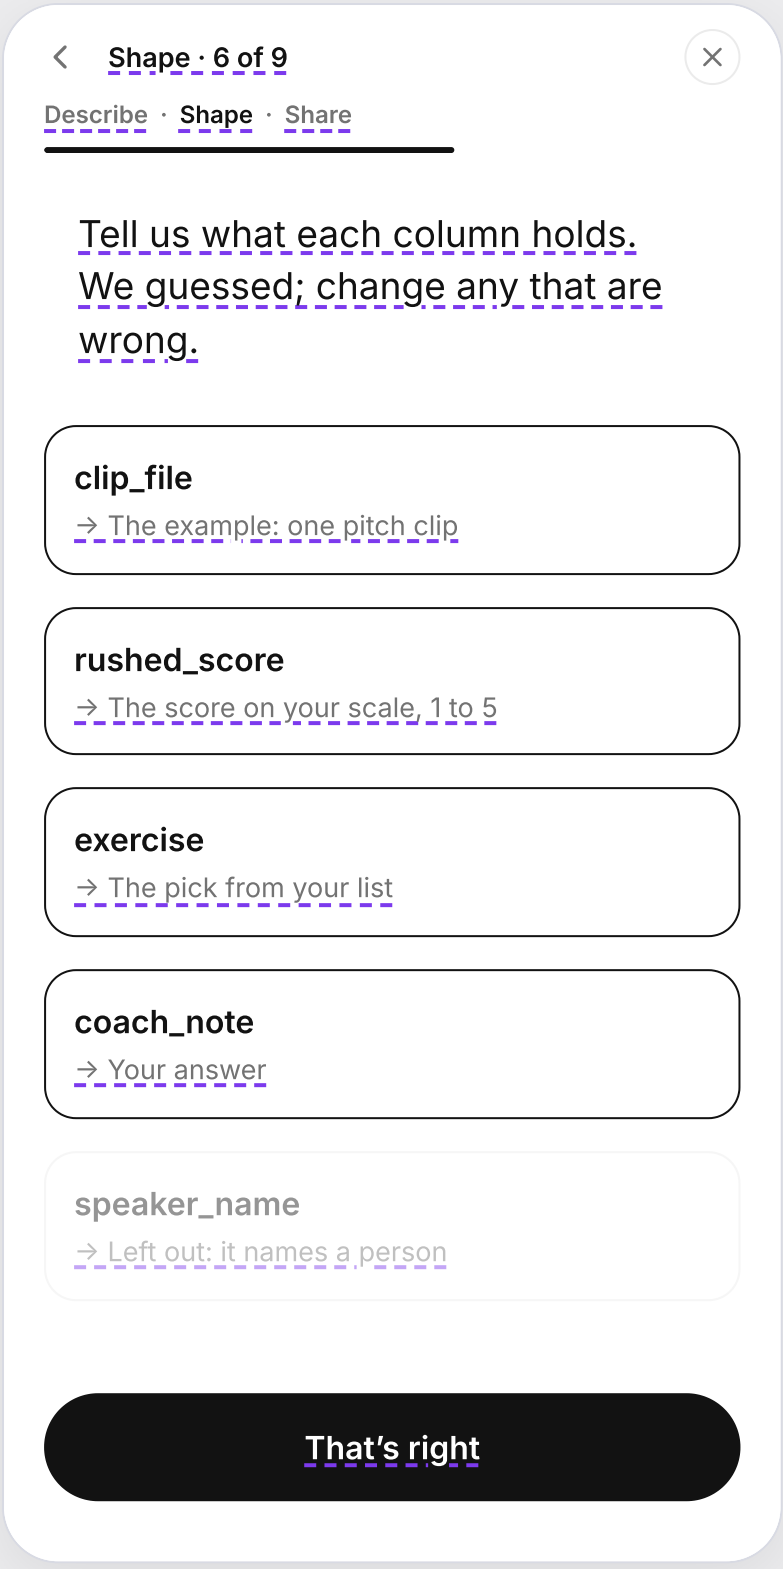
7/18 Changed, from the ML-engineering review. GPT-4o mini reads and writes text; it cannot hear a voice. So a score on how someone sounds needs a small model that learns from the voice itself, the way Speaking confidence does. The app picks the right AI for each part, and GPT-4o mini stays the one for words.



8/18 Changed, from the data-science review. Each rule now shows who it was measured on, and the app tests every rule against the teachers' first ratings before trusting it. (Sample file and numbers.)



9/18 New. The setup accepts a dataset the person already has, so a model need not start from nothing. (Sample files.)



10/18 The app maps the file onto the model's three parts. A column that names a person is left out by default.

Shape · 7 of 9

Describe · Shape · Share

We checked your data.

Rows we can use	**1,180 / 1,240**
Repeats removed	**38**
Near repeats removed	**17**
Rows with no score, skipped	**22**
Rows that look mislabelled	**42 · checked first**
Names and places removed	**214**



Most of your examples sit at 2 or 3. Your model will learn little about racing pitches (5) until your

Upload it

Fix the file first

11/18 New, with both reviews in mind. The check a data scientist would run first: what is usable, repeats, missing scores, how the scores spread across the scale (and where the model will be blind), and names removed. Uploaded labels come from unknown raters, so teachers re-rate a sample before they count, exactly as the starting rules are checked. The final exam never uses it. (Sample numbers.)

Shape · 8 of 9

Describe · Shape · Share

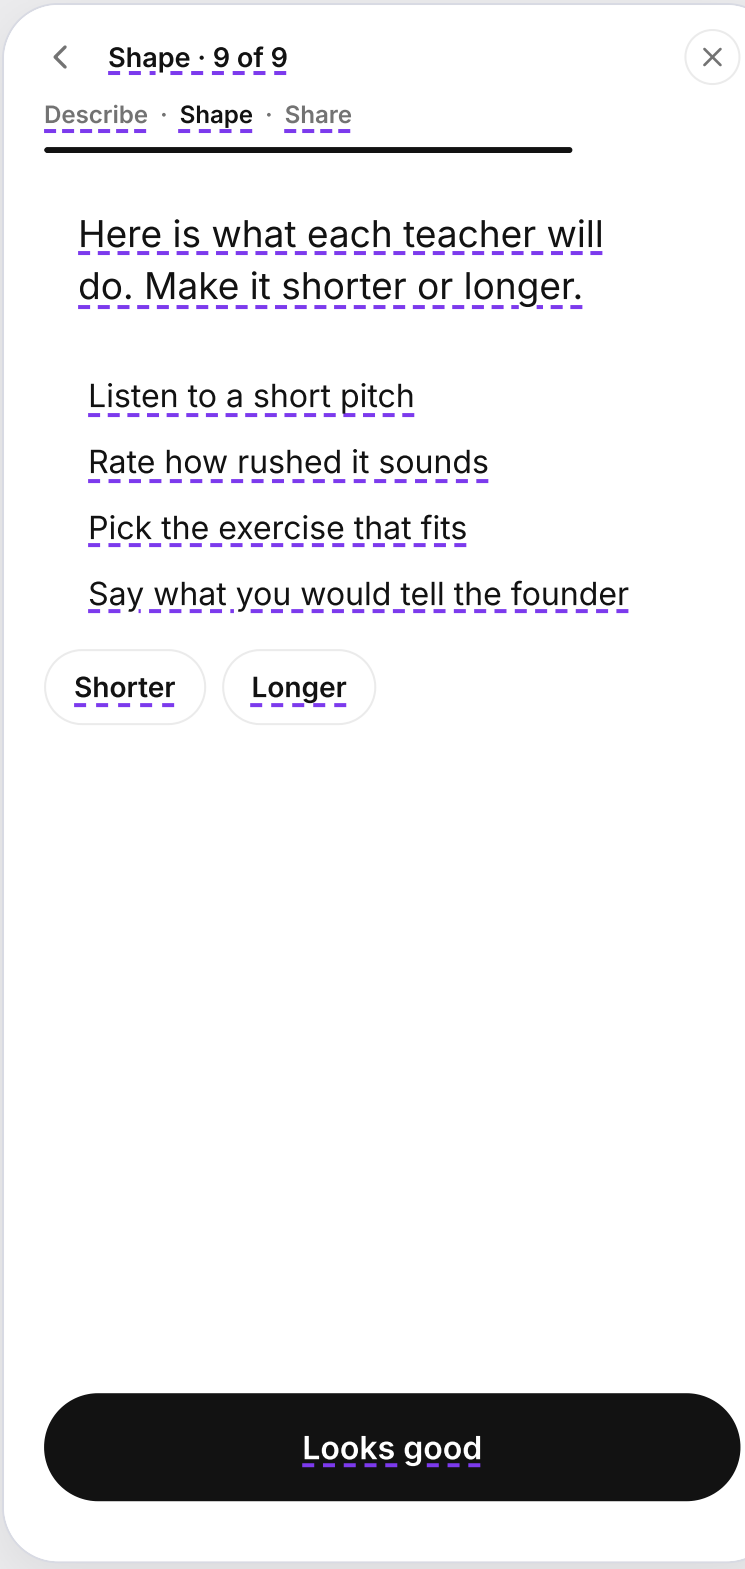
Before anyone teaches it, two experts set aside a final exam: 100 clips your model never learns from.

It can be sold only when it does clearly better on this exam than it does today.

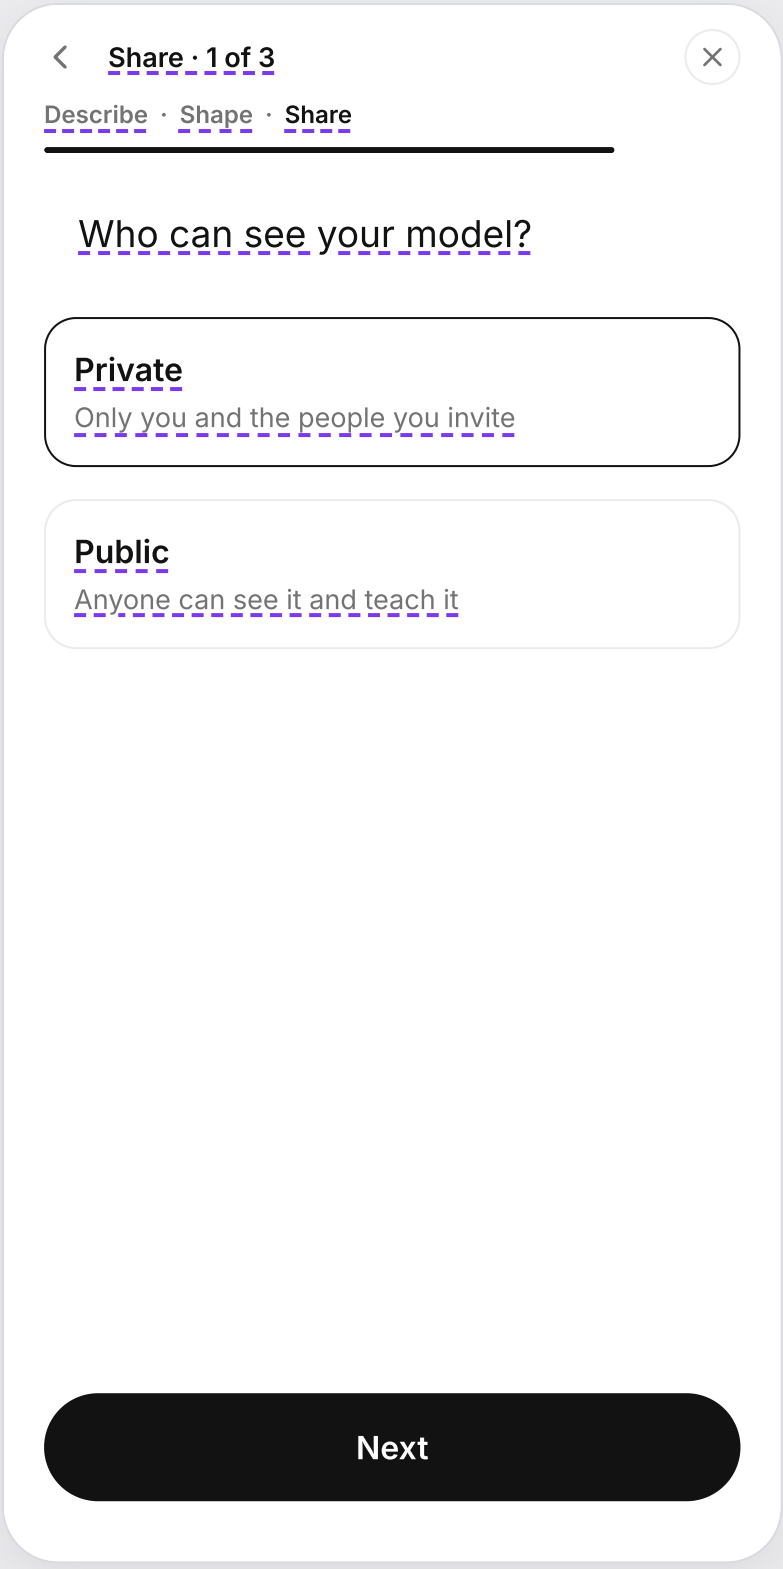
- [x] **You**
  Owner
- [x] **Daniel Okafor**
  Speaking coach

Ask Daniel

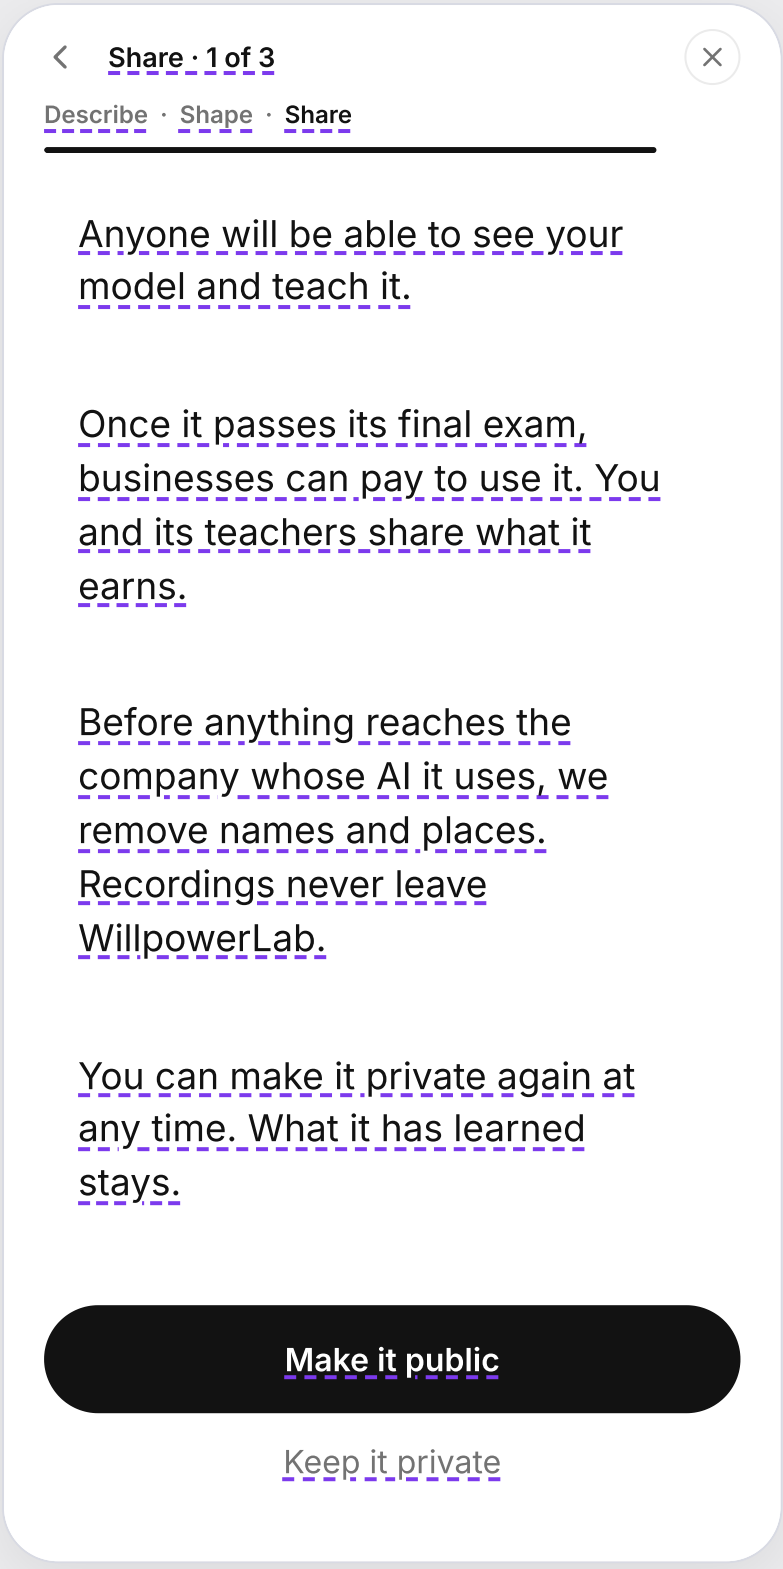
12/18 New, from both reviews. The exam is written first and kept sealed, by two experts whose agreement is measured. Fifty cases can only show a big improvement; 100 shows a real one. The bar is to beat the model as it works today on its starting rules, not merely to exist.



13/18 The journey the app generated. Shorter or longer switches parts off or on.



14/18 One switch. Tap **Public** to see its disclaimer.



15/18 The disclaimer before going public, now also saying what reaches the AI company.

Share · 2 of 3

Describe · Shape · Share

Who should teach it with you?

- [ ] **Everyone in Northwind sales team**
  14 people · each one chooses whether to join
- [x] **Maya Lindqvist**
  Head of Sales
- [x] **Daniel Okafor**
  Speaking coach
- [ ] **Amira Haddad**
  Account executive

Invite them

Copy an invite link

16/18 Invite people from your chats, or copy a link.

Share · 3 of 3

Describe · Shape · Share

Pitch Coach

0%

Pitch Coach is ready to learn. Today it answers from your starting rules and examples.

To finish setting it up

- [x] ~~The scale~~
- [x] ~~What it will do~~
- [ ] Starting rules · optional
- [ ] Your existing data · optional
- [ ] The final exam
- [ ] Invite teachers

Sell it through an API

This opens once its first training beats today's version on the final exam. It needs about 500 ratings from trusted teachers first.

Try it

Back to the list

17/18 Changed, from the UX review: a model can start right after its scale ("Start with this, finish the rest later" on the scale screen); what is left becomes a checklist on its page. The API opens after the first training beats today's version.

Skip ahead two weeks ›

Pitch Coach

Pitch Coach

Needs you now

Your teachers disagree most between 3 (Hurried in places) and 4 (Rushed). Clearer words for those two points would help.

Make those two clearer

36% Pitch Coach is 36% of the way to its first training: about 4 to 6 weeks more. The slowest jar decides: different voices.

How often your teachers agree	**Within one point 9 times in 10 · exactly 6 in 10**
Ratings from trusted teachers	**180 / 500**
Final exam	**100 of 100 · sealed**
Your uploaded examples	**1,180 · they count: teachers agreed with 8 in 10**

Where its examples come from

Back to the list

18/18 Changed, from both reviews: one card on top, **Needs you now**; agreement shown within one point and exactly; the details folded (tap to open). What the owner needs to steer. **Earlier, from the data-science review:** an honest range for when it will be ready and which jar decides it; where the teachers disagree on the scale, with the fix; proof the exam is fair (no speaker on both sides, every point covered); and the gaps, which the teachers already hear first. Totals only, never a person.

Job interview coach

Job interview coach

100% Job interview coach passed its final exam. Nine interview coaches taught it.

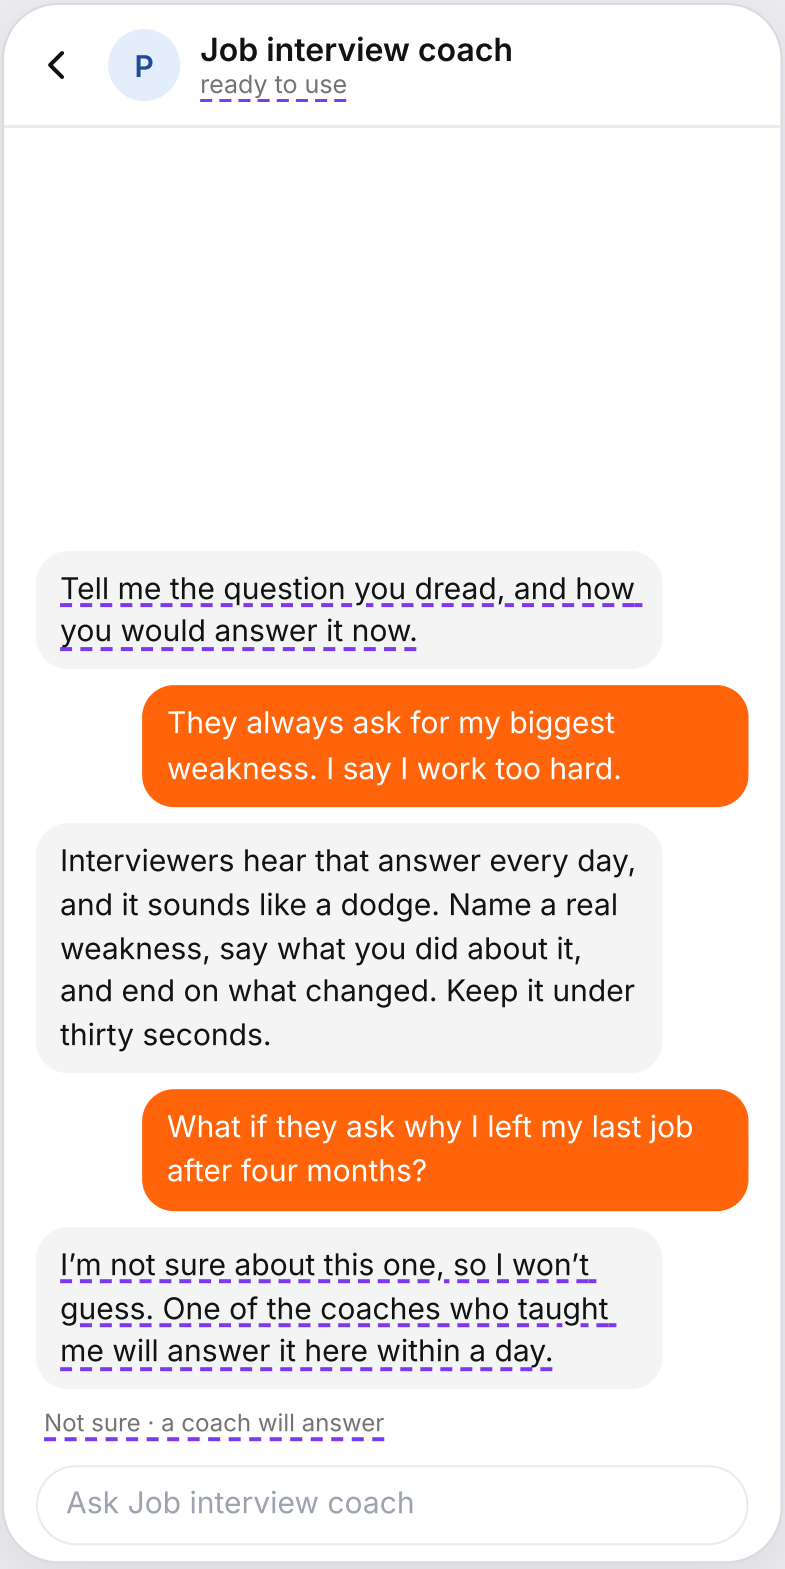
It helps people prepare for job interviews: it hears the question they fear, and answers the way an experienced interview coach would.

Good at	**Common interview questions, in English**
Not so good at	**Technical interviews; other languages**
Not for	**Legal or medical advice**
On its final exam	**Coaches who didn't know which answer was whose rated its answers as good as a coach's about 8 times in 10.**
When it isn't sure	**It says so, and a coach answers instead**

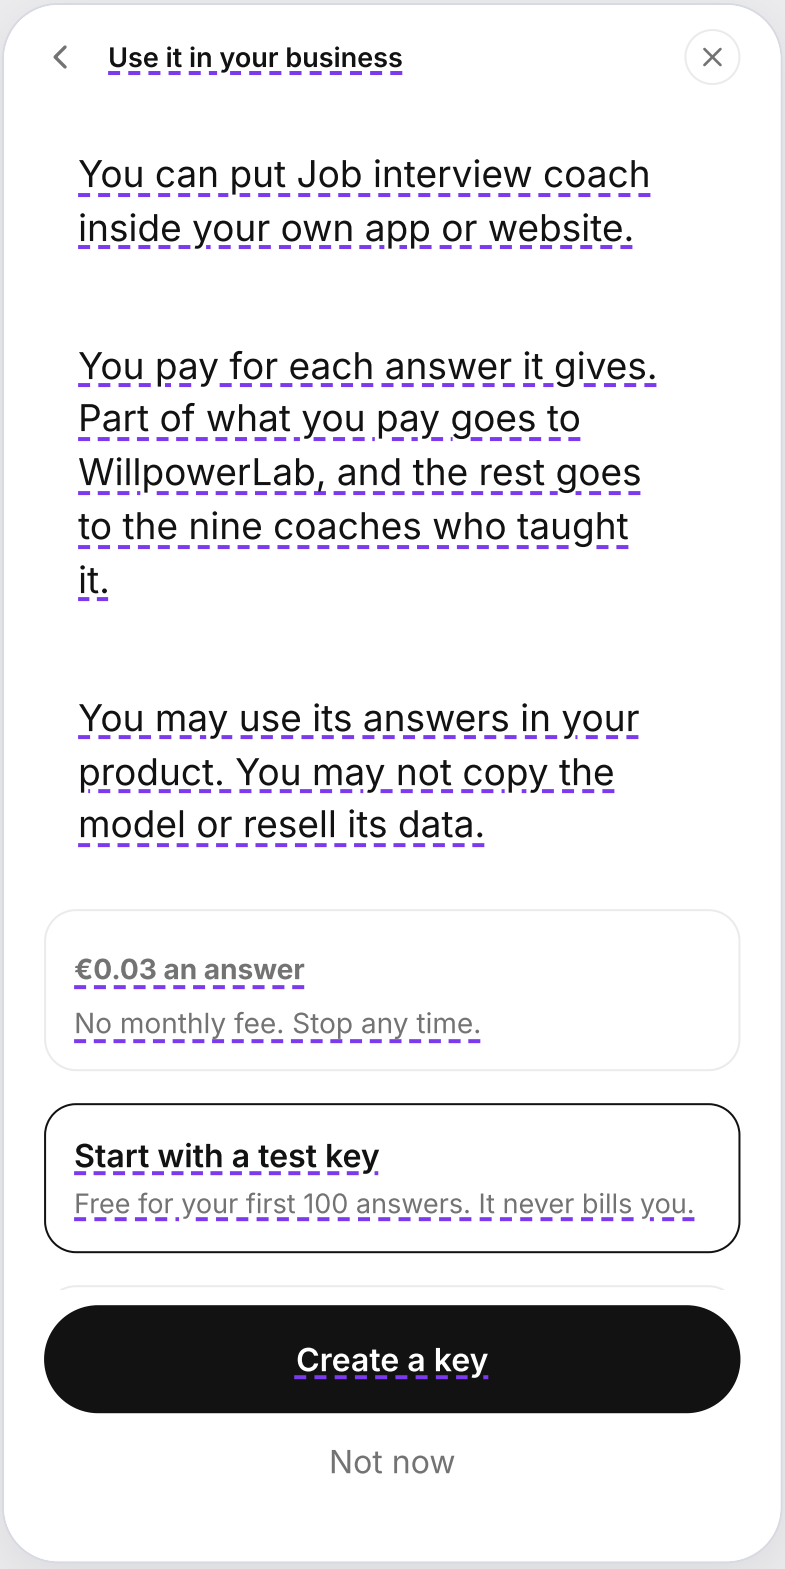
Use it in your business

Try it

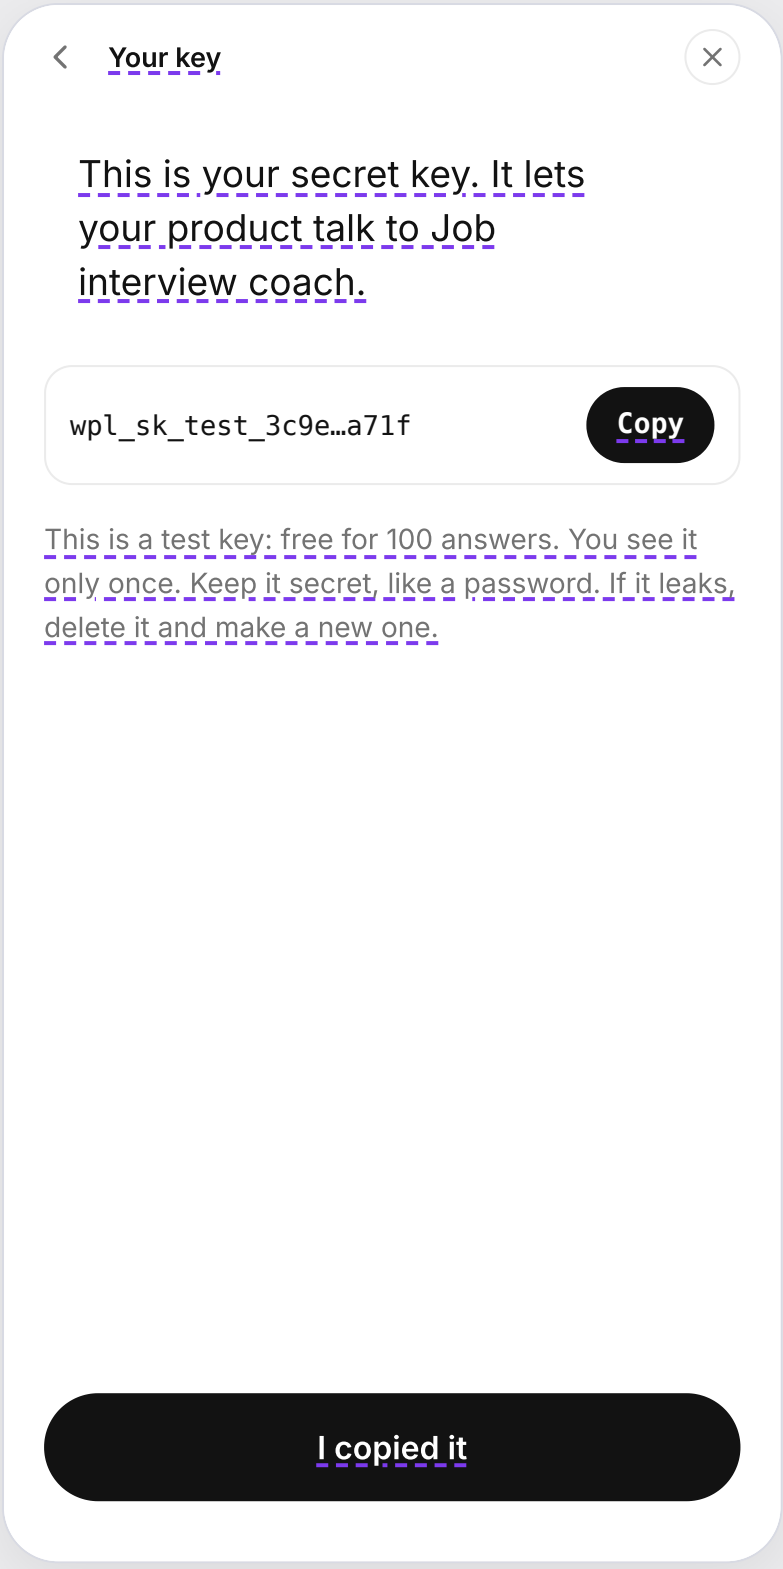
1/6 A business owner (say, a careers website) finds a model that is ready to use. **New, from the data-science review:** what it is good at, where it struggles, what it is not for, and its version. Tap **Try it** first.



2/6 Trying it on its link, in the same chat box as Will's, from the app's code. **New, from the data-science review:** when it isn't sure, it says so instead of guessing, and a coach answers. When it convinces you, go back and tap **Use it in your business**.



3/6 Before the key: what it costs, who gets paid, and what the business may do. **Changed, from the ML-engineering review:** a free test key first, a monthly limit, and a promise about the business's users' data. (Sample price.)



4/6 The key, shown once, as at any AI company. Tap **Copy**. (Sample key.)

Put it in your product

Copied

Paste the key into your product, then send a question like this one. The answer comes back as text.

JavaScript | Python | curl

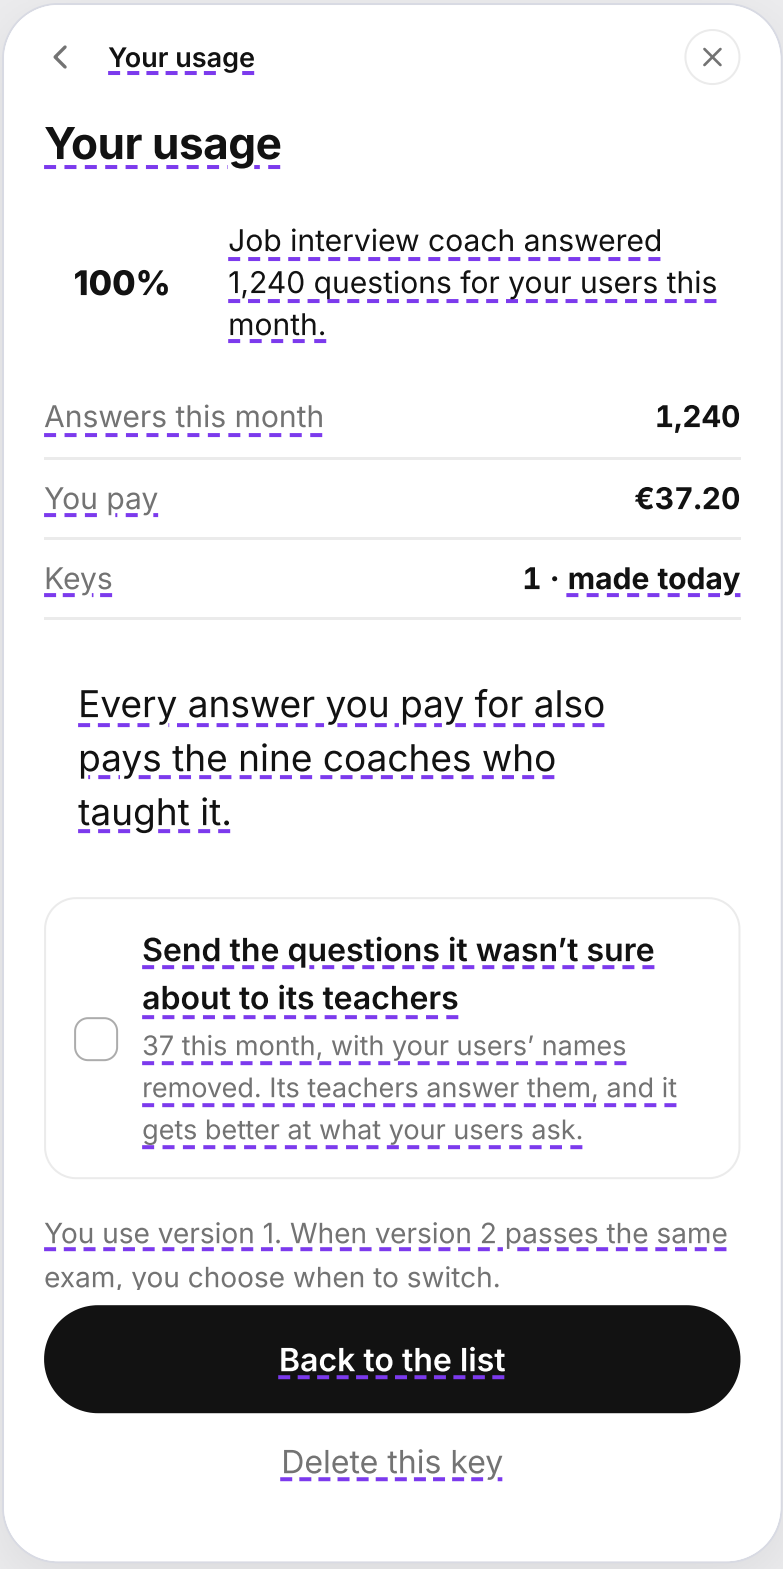
```
const res = await
fetch("https://api.willpowerlab.com/v1/model
s/job-interview-coach/answer", {
  method: "POST",
  headers: { Authorization: `Bearer
${process.env.WILLPOWERLAB_KEY}` },
  body: JSON.stringify({ message: "How do I
answer: what is your biggest weakness?" }),
});
const { answer, sure } = await res.json();
```

Each answer comes with “sure”: true or false. When it is false, show the answer as a draft, or pass the question to a person.

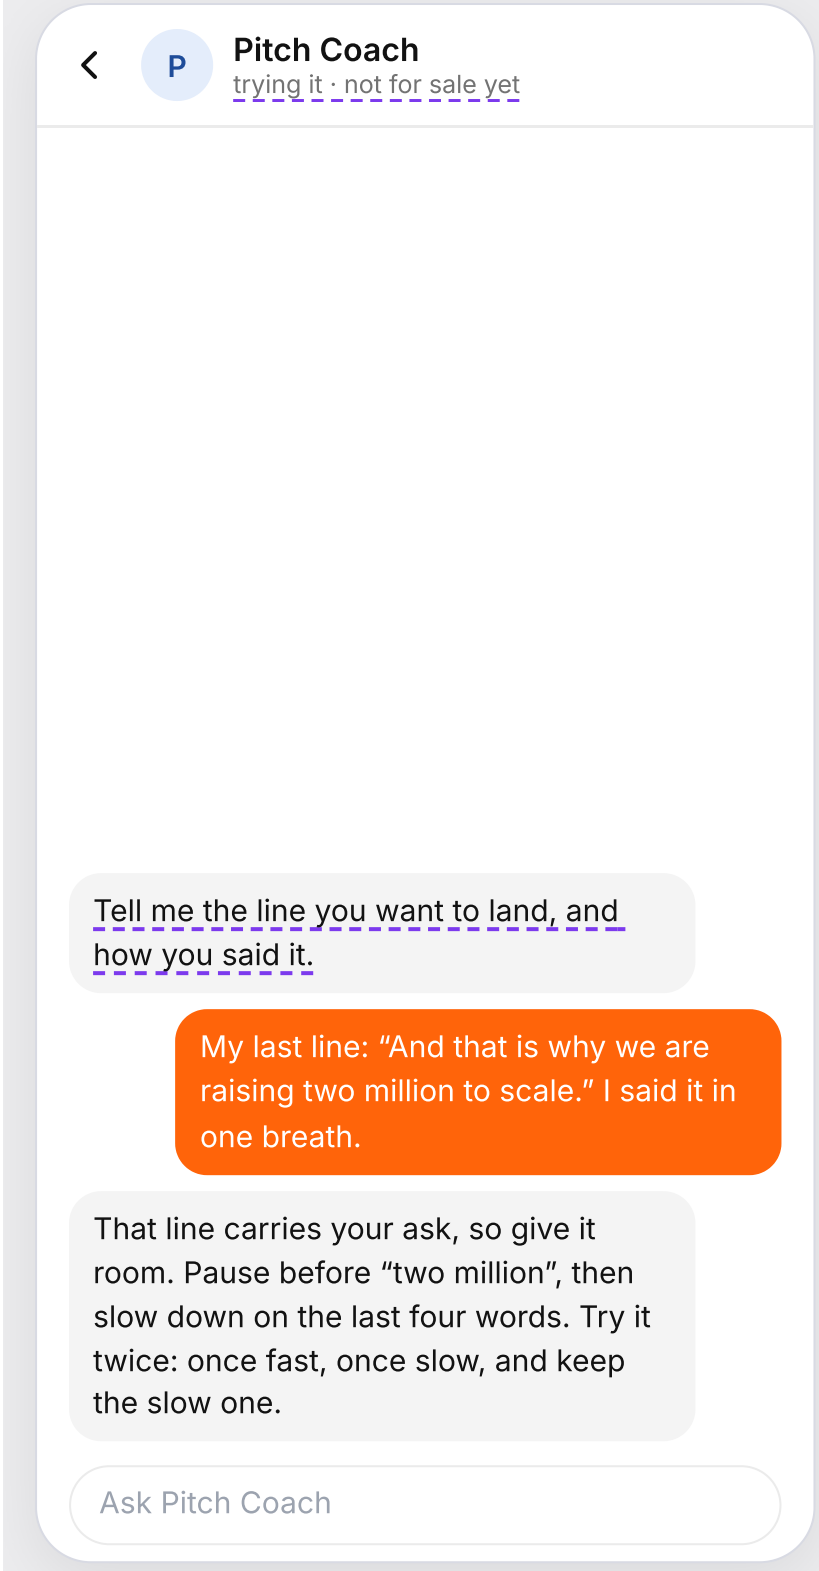
Done

Send it to your developer

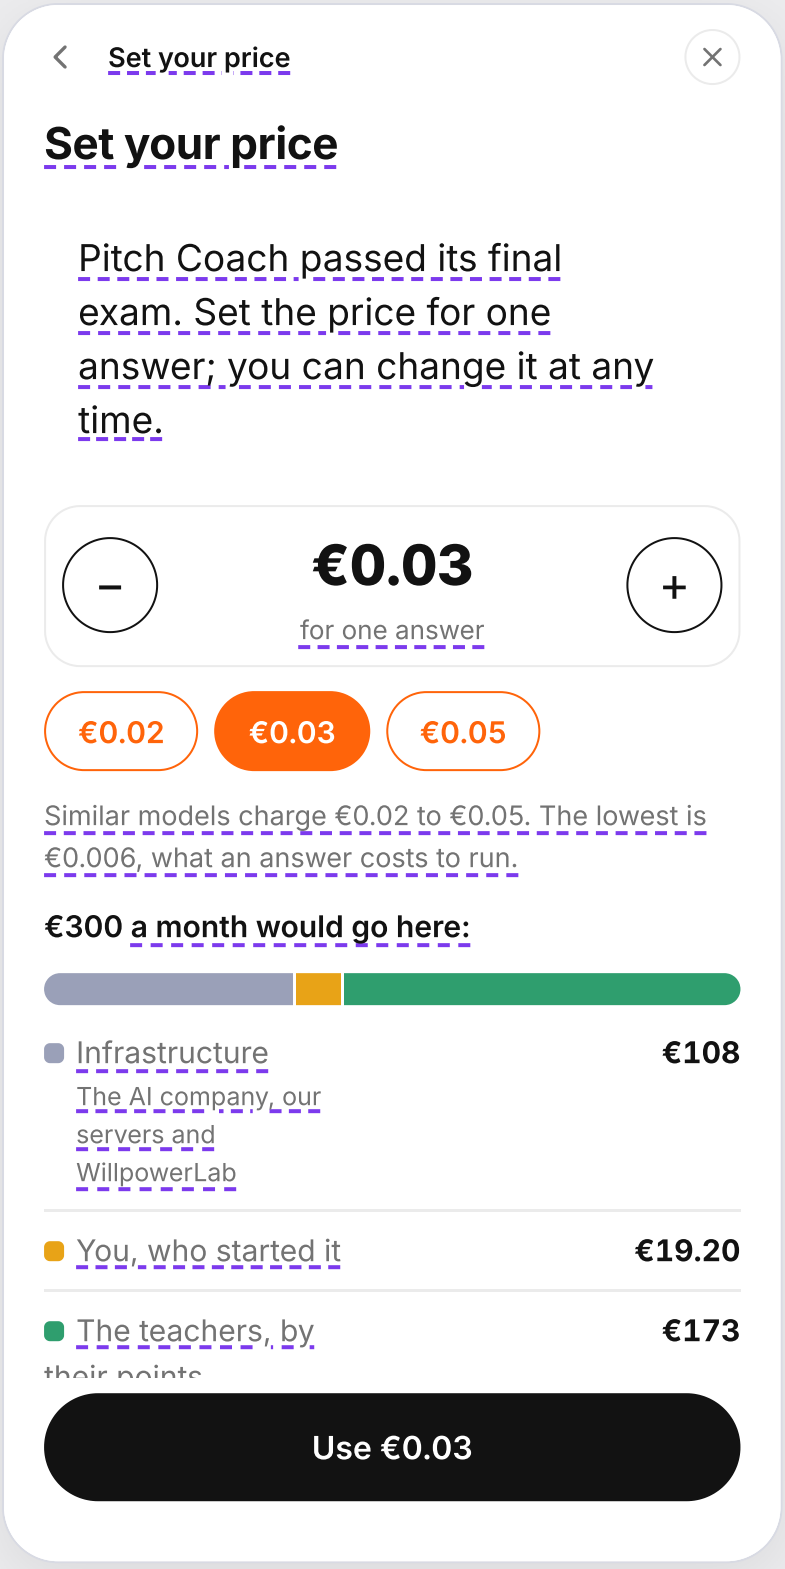
5/6 One request, ready to copy. **Changed, from both reviews:** tabs for the language the developer uses, “Send it to your developer”, and a “sure” field in every answer, so a product can treat an unsure answer as a draft. (The address is an example.)



6/6 What the business sees afterwards. **New, from both reviews:** the questions the model wasn't sure about can go back to its teachers, with the business's yes, which pays them and fills the model's gaps; and versions are pinned, so an update never changes a business's product without its choice. (Sample numbers.)



1/4 Trying a model on its link: a chat with an expert in one field. Text first.



2/4 Try it: change the price with − and +, or tap a suggestion, and move the slider. The split follows. **Infrastructure** is one row now: the running costs and WillpowerLab's part together. WillpowerLab's part is taken from what is left after running costs, so at the lowest price everything goes to infrastructure. (Sample fee of 20% and founder's share of 10%: both are yours to set.)

Your earnings

Your earnings

100% Pitch Coach version 1 earned you €12.00 this month.

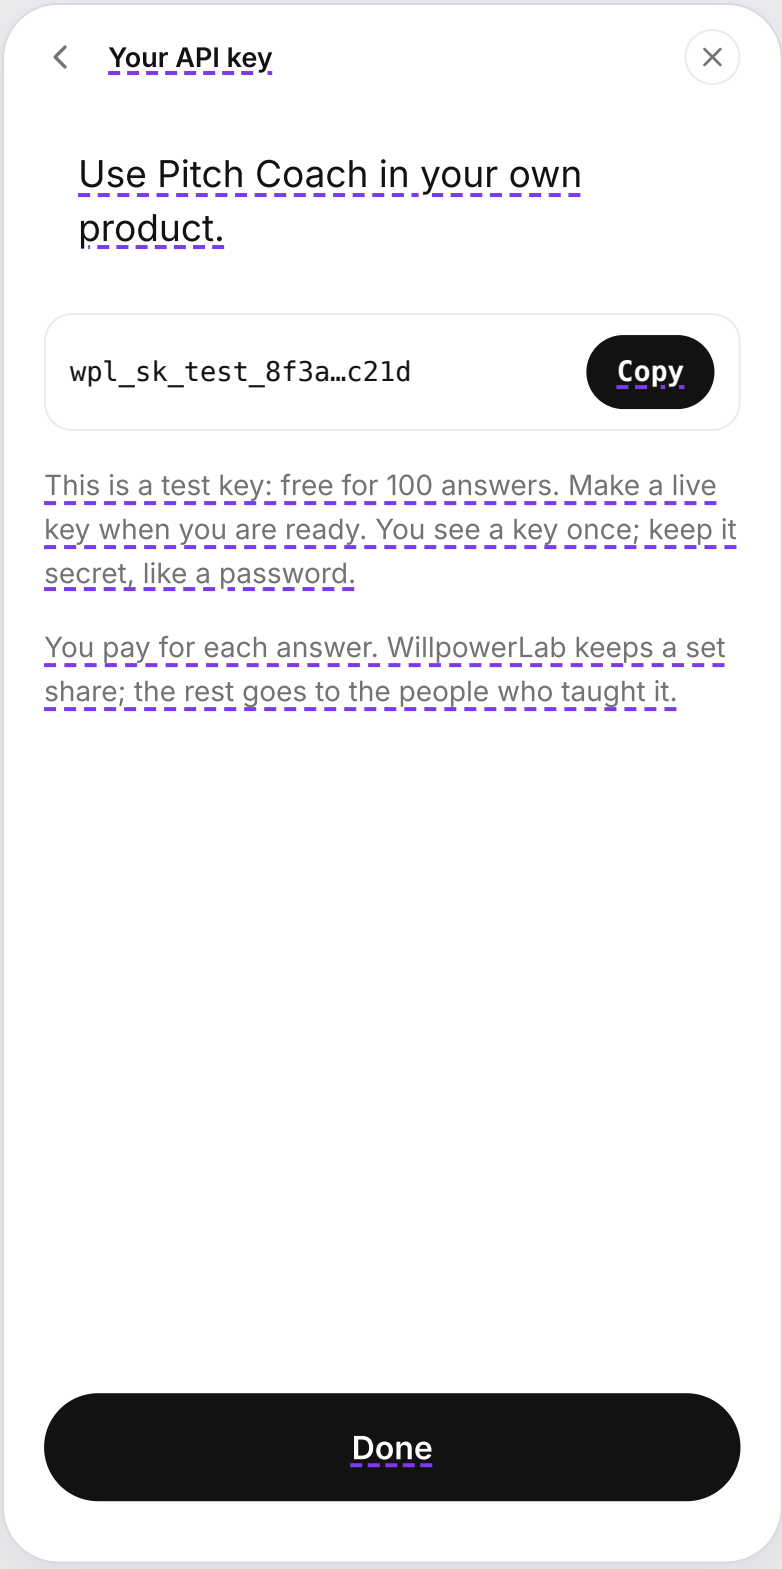
Your share of version 1	**7%**
Your answers that trained it	**214**
Next payout	**1 December · to your bank**

Your share is fixed for version 1. It grows with each answer that counts; hard clips count for more. Answers you give now count toward version 2.

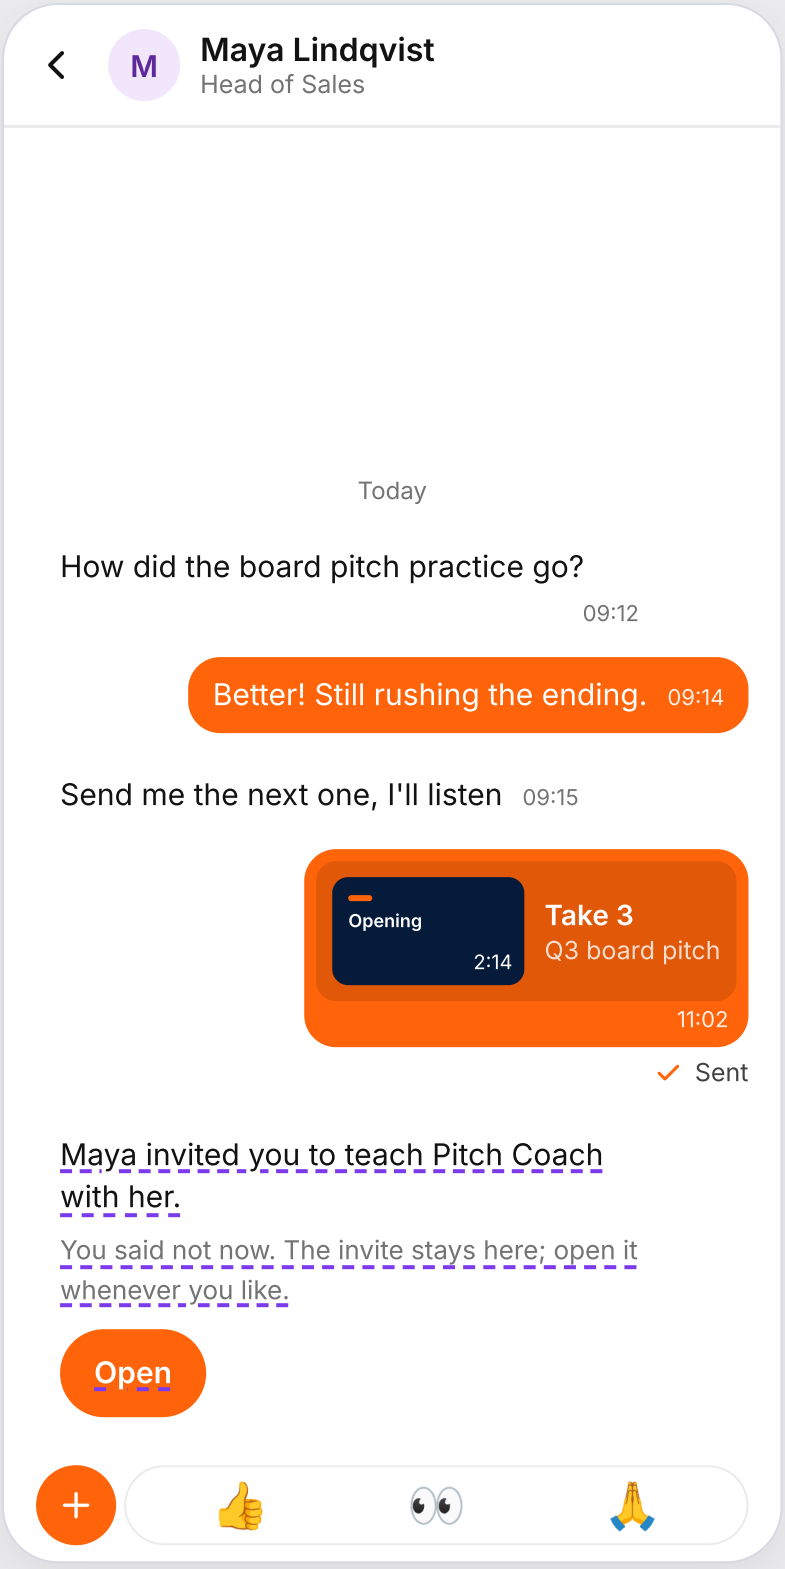
We pay once a month when you have earned €20 or more. Before your first payout, we ask for your name and bank details, as every payment service must.

Back to the list

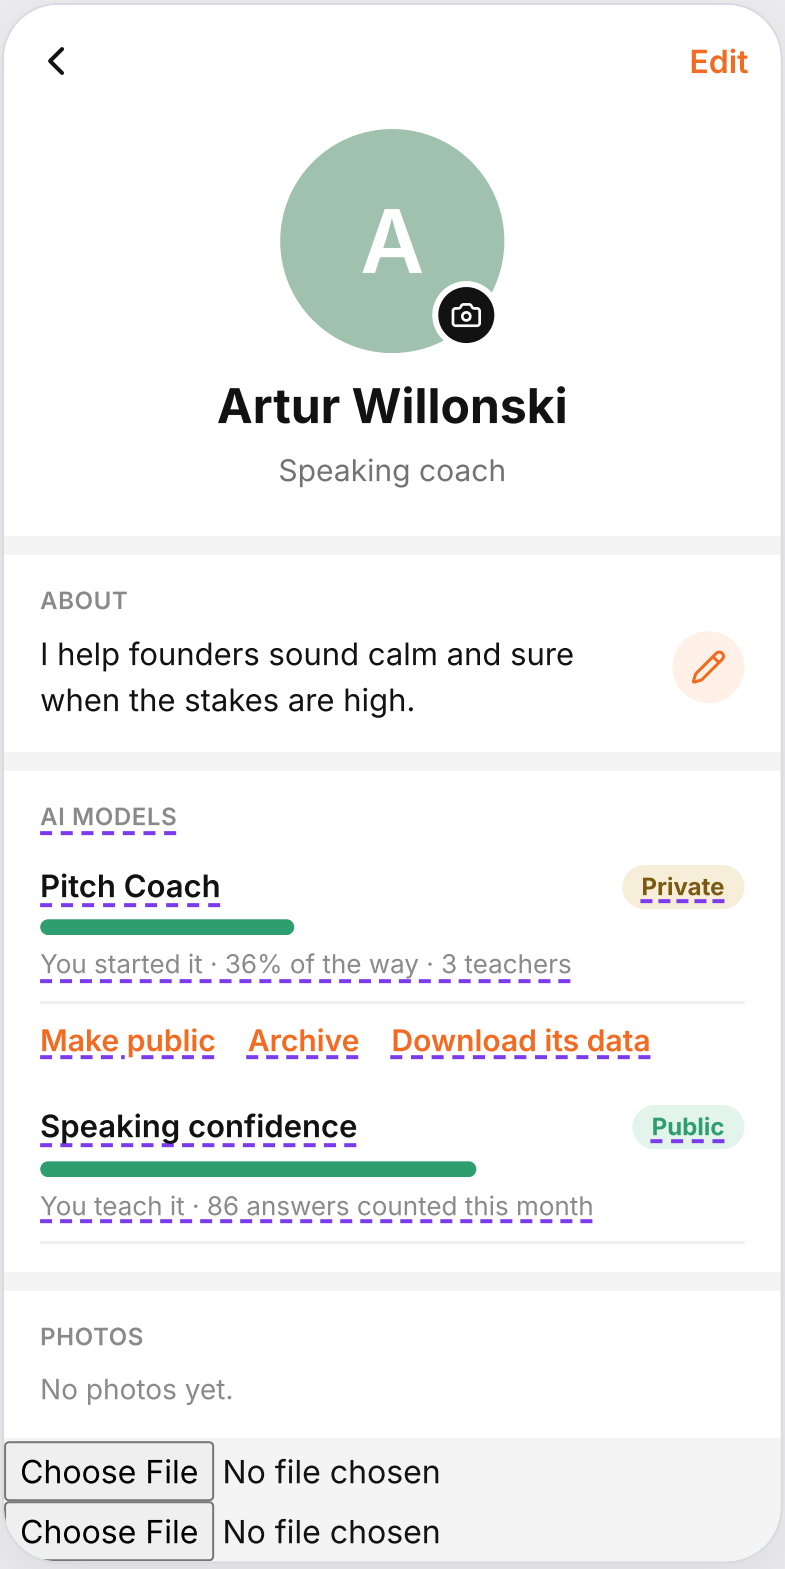
3/4 New, from N531. What a teacher sees: their share of the version their answers trained, what it earned, and when it pays out. Money, monthly, above a minimum, once the legal terms are signed; tokens until then. (Sample numbers.)



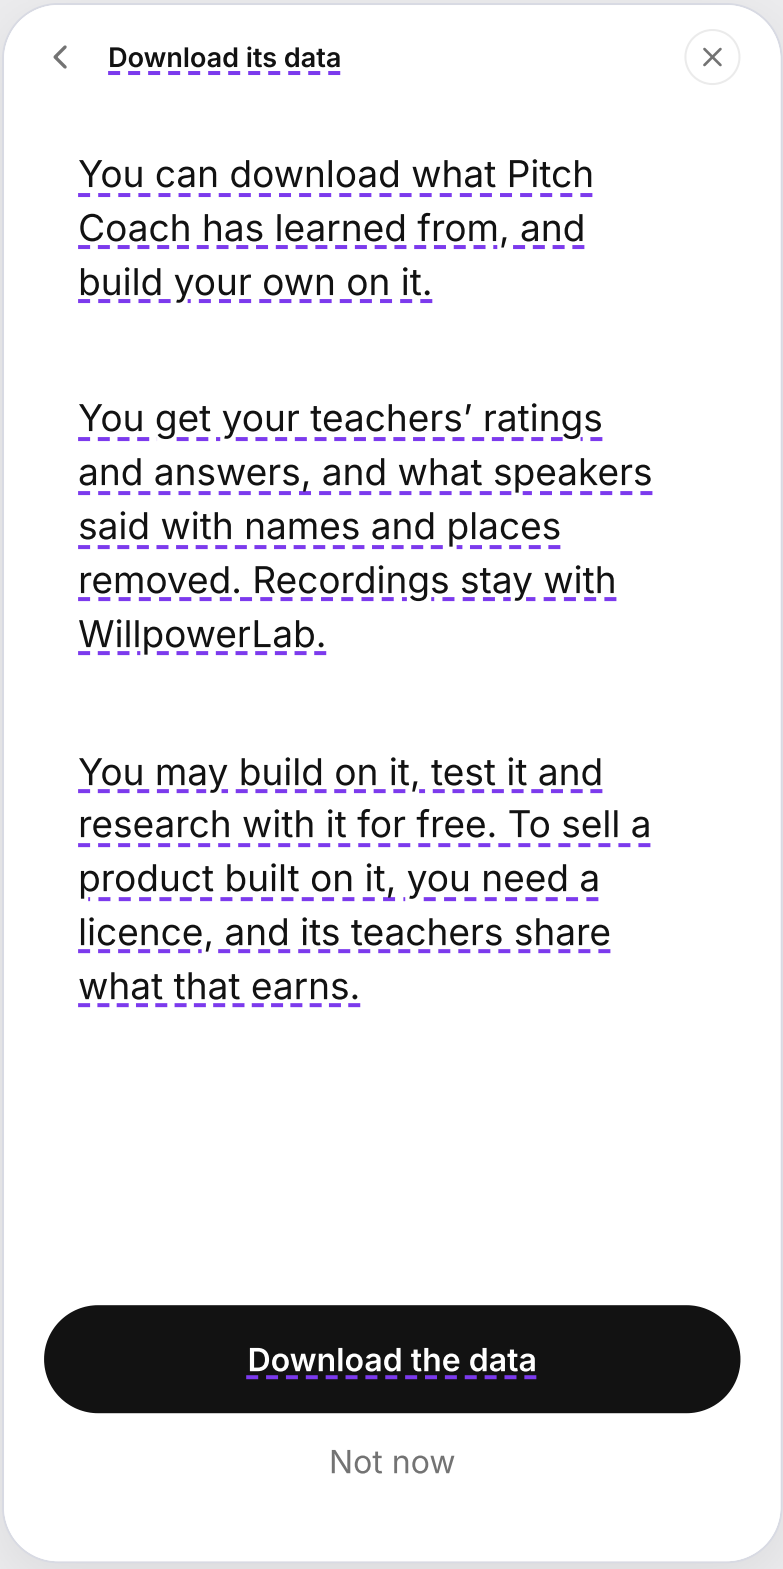
4/4 After the first training passes: create a secret key, as at any AI company. (Sample key.)



1/3 An invite to teach stays in the chat after “Not now”. Scroll up, tap Open.



2/3 Your real profile screen, from the app's code, with one new section in its own style: **AI models**, the ones you started (make public, archive, download its data) and the ones you teach. Tap **Download its data**.



3/3 Open on WillpowerLab, not open source: the data is free to download for owners and contributors; using the model through its API is paid, and selling something built on the data needs a licence. Speakers' words are included only once the training yes says so.